

RelaxIV Documentation

version 1.80

date 11 - 02 - 2011

Operations Research Group

Dipartimento di Informatica

Università di Pisa

Chapter 1

Manual

1.1 Introduction

This file is a short user manual for our distribution of RelaxIV, a C++ implementation of the latest variants of Relaxation algorithms for linear Min Cost Flow problems.

1.2 Standard Disclaimer

This code is provided "as is", without any explicit or implicit warranty that it will properly behave or it will suit you needs. Although codes reaching the distribution phase have usually been extensively tested, we cannot guarantee that they are absolutely bug-free (who can?). Any use of the codes is at you own risk: in no case we could be considered liable for any damage or loss you would eventually suffer, either directly or indirectly, for having used this code. More details about the non-warranty attached to this code are available in the license description file.

The code also comes with a "good will only" support: feel free to contact us for any comments/critics/bug report/request help you may have, we will be happy to try to answer and help you. But we cannot spend much time solving your problems, just the time to read a couple of e-mails and send you fast suggestions if the problem is easily solvable. Apart from that, we can't offer you any support.

1.3 License

This code is provided free of charge for academic purposes under the "academic license": see the file academici.txt. Commercial use of this code is only allowed at a fee: please contact the original author dimitrib@mit.edu for further information.

1.4 Release

Current version of RelaxIV is: 1.80 (after a major renumbering) Current date is: February 11, 2011

1.5 How to Use It

This release comes out with the following files:

- Manual This file.
- doc/academicl.txt Description of the "academic license".
- doc/refman.pdf Pdf version of the full manual.
- doc/html/* Html version of the full manual.
- OPTUtils.h Contains some small utilities such as random number generators and timing routines.
- MCFCClass.h Contains the declaration of class MCFCClass; it is an "abstract" class with "pure virtual" methods, so that you cannot declare objects of type MCFCClass, but only of derived class of its. MCFCClass offers a general interface for MCF solver codes, with methods for setting and reading the data of the problems, for solving it and retrieving solution informations, and so on. The actual Min Cost Flow solver distributed in this package conforms with the interface, i.e., it derives from MCFCClass; however, the idea is that applications using this interface would require almost no changes if any other of the available solvers implementing the interface is used. Carefully read the public interface of the class to understand how to use the public methods of the class. Also, some compile time switches let you choose important features of the algorithm, such as if the node "names" are the integers $0 \dots n - 1$ rather than $1 \dots n$.
- RelaxIV.h Contains the declarations of class RelaxIV. Being a derived class of MCFCClass, most of its interface is defined and discussed in MCFCClass.h; however, some implementation dependent details (and compile-time switches) which are worth knowing are described in this file.
- RelaxIV.C Contains the implementation of the RelaxIV class, which uses the Relaxation algorithm version IV (due to Dimitri P. Bertsekas) to solve the Min Cost Flow problem. You should not need to read it.

The following files are not a part of the actual MCF solver, they are only provided as an example of how to use it:

- Main.C Contains an example of use of RelaxIV; it reads the data from a file in DIMACS standard format, constructs the problem, solves it and prints the results.
- Config.txt Allows (if exists) to set some important parameters for the solver.
- sample.dmx An example of a tiny (4 nodes, 4 arcs) network in DIMACS standard format; the comments in the file also describe the format itself.
- makefile A makefile of the example.

1.6 Testing the distribution

To compile the example, edit the makefile and do the proper changes (such as the compiler name), then just type "make". To execute the example

MCFSolve sample.dmx

The output should be something like Optimal Objective Function value = 4 Solution time (s): user 0, system 0

Contents

1	Manual	i
1.1	Introduction	i
1.2	Standard Disclaimer	i
1.3	License	i
1.4	Release	i
1.5	How to Use It	ii
1.6	Testing the distribution	ii
2	Module Index	1
2.1	Modules	1
3	Class Index	3
3.1	Class Hierarchy	3
4	Class Index	5
4.1	Class List	5
5	File Index	7
5.1	File List	7
6	Module Documentation	9
6.1	Compile-time switches in MCFCClass.h	9
6.1.1	Detailed Description	9
6.1.2	Define Documentation	9
6.1.2.1	USERNAME0	9
6.2	Classes in MCFCClass.h	10
6.3	Compile-time switches in OPTUtils.h	11
6.3.1	Detailed Description	11
6.3.2	Define Documentation	11
6.3.2.1	OPT_RANDOM	11

6.3.2.2	OPT_TIMERS	11
6.3.2.3	OPT_USE_NAMESPACES	12
6.4	Classes in OPTUtils.h	13
6.5	Functions in OPTUtils.h	14
6.5.1	Function Documentation	14
6.5.1.1	DfltSfInpt	14
6.6	Compile-time switches in RelaxIV.h	15
6.6.1	Detailed Description	15
6.6.2	Define Documentation	15
6.6.2.1	AUCTION	15
6.6.2.2	DYNMC_MCF_RIV	15
6.6.2.3	RELAXIV_STATISTICS	15
6.6.2.4	SAME_GRP_H_RIV	16
6.7	Classes in RelaxIV.h	17
7	Class Documentation	19
7.1	MCFCClass::Eps< T > Class Template Reference	19
7.1.1	Detailed Description	19
7.2	MCFCClass::Inf< T > Class Template Reference	20
7.2.1	Detailed Description	20
7.3	MCFCClass Class Reference	21
7.3.1	Detailed Description	27
7.3.2	Member Typedef Documentation	28
7.3.2.1	FONumber	28
7.3.3	Member Enumeration Documentation	28
7.3.3.1	MCFAnswer	28
7.3.3.2	MCFFFFmt	28
7.3.3.3	MCFParam	28
7.3.3.4	MCFStatus	29
7.3.4	Constructor & Destructor Documentation	29
7.3.4.1	MCFCClass	29
7.3.4.2	~MCFCClass	29
7.3.5	Member Function Documentation	30
7.3.5.1	AddArc	30
7.3.5.2	AddNode	30
7.3.5.3	ChangeArc	30
7.3.5.4	CheckDSol	30

7.3.5.5	CheckPSol	30
7.3.5.6	ChgCost	30
7.3.5.7	ChgCosts	31
7.3.5.8	ChgDfct	31
7.3.5.9	ChgDfcts	31
7.3.5.10	ChgQCoef	32
7.3.5.11	ChgQCoef	32
7.3.5.12	ChgUCap	32
7.3.5.13	ChgUCaps	33
7.3.5.14	CloseArc	33
7.3.5.15	DelArc	33
7.3.5.16	DelNode	33
7.3.5.17	ETZ	33
7.3.5.18	GetPar	34
7.3.5.19	GetPar	34
7.3.5.20	GEZ	34
7.3.5.21	GT	34
7.3.5.22	GTZ	34
7.3.5.23	HaveNewPi	34
7.3.5.24	HaveNewX	35
7.3.5.25	IsClosedArc	35
7.3.5.26	IsDeletedArc	35
7.3.5.27	LEZ	35
7.3.5.28	LoadDMX	35
7.3.5.29	LoadNet	36
7.3.5.30	LT	37
7.3.5.31	LTZ	37
7.3.5.32	MCFArcs	37
7.3.5.33	MCFCost	37
7.3.5.34	MCFCosts	37
7.3.5.35	MCFCosts	38
7.3.5.36	MCFDfct	38
7.3.5.37	MCFDfcts	38
7.3.5.38	MCFDfcts	38
7.3.5.39	MCFENde	39
7.3.5.40	MCFENdes	39

7.3.5.41	MCFGetDFO	39
7.3.5.42	MCFGetFO	39
7.3.5.43	MCFGetPi	39
7.3.5.44	MCFGetPi	40
7.3.5.45	MCFGetRC	40
7.3.5.46	MCFGetRC	40
7.3.5.47	MCFGetRC	40
7.3.5.48	MCFGetState	41
7.3.5.49	MCFGetStatus	41
7.3.5.50	MCFGetUnbCycl	41
7.3.5.51	MCFGetUnfCut	42
7.3.5.52	MCFGetX	42
7.3.5.53	MCFGetX	42
7.3.5.54	MCFm	43
7.3.5.55	MCFmmax	43
7.3.5.56	MCFn	43
7.3.5.57	MCFnmax	43
7.3.5.58	MCFPutState	43
7.3.5.59	MCFQCoef	44
7.3.5.60	MCFQCoef	44
7.3.5.61	MCFQCoef	44
7.3.5.62	MCFSNde	44
7.3.5.63	MCFSNdes	44
7.3.5.64	MCFUCap	45
7.3.5.65	MCFUCaps	45
7.3.5.66	MCFUCaps	45
7.3.5.67	OpenArc	45
7.3.5.68	PreProcess	45
7.3.5.69	SetMCFTIME	45
7.3.5.70	SetPar	46
7.3.5.71	SetPar	46
7.3.5.72	SolveMCF	47
7.3.5.73	TimeMCF	47
7.3.5.74	TimeMCF	47
7.3.5.75	WriteMCF	47
7.3.6	Member Data Documentation	48

7.3.6.1	MaxIter	48
7.3.6.2	MaxTime	48
7.3.6.3	status	48
7.4	MCFClass::MCFException Class Reference	49
7.4.1	Detailed Description	49
7.5	MCFClass::MCFState Class Reference	50
7.5.1	Detailed Description	50
7.6	OPTrand Class Reference	51
7.6.1	Detailed Description	51
7.6.2	Member Function Documentation	51
7.6.2.1	rand	51
7.7	RelaxIV Class Reference	52
7.7.1	Detailed Description	53
7.7.2	Member Enumeration Documentation	54
7.7.2.1	MCFRParam	54
7.7.2.2	RIVFIFmt	54
7.7.3	Constructor & Destructor Documentation	54
7.7.3.1	RelaxIV	54
7.7.4	Member Function Documentation	54
7.7.4.1	AddArc	54
7.7.4.2	AddNode	54
7.7.4.3	ChangeArc	54
7.7.4.4	CloseArc	55
7.7.4.5	DelArc	55
7.7.4.6	DelNode	55
7.7.4.7	GetPar	55
7.7.4.8	IsClosedArc	55
7.7.4.9	IsDeletedArc	55
7.7.4.10	LoadNet	56
7.7.4.11	MCFCost	56
7.7.4.12	MCFCosts	56
7.7.4.13	MCFDfct	56
7.7.4.14	MCFDfcts	57
7.7.4.15	MCFENde	57
7.7.4.16	MCFENdes	57
7.7.4.17	MCFGetFO	57

7.7.4.18	MCFGetPi	58
7.7.4.19	MCFGetPi	58
7.7.4.20	MCFGetRC	58
7.7.4.21	MCFGetRC	58
7.7.4.22	MCFGetState	58
7.7.4.23	MCFGetX	59
7.7.4.24	MCFPutState	59
7.7.4.25	MCFSNde	59
7.7.4.26	MCFSNdes	60
7.7.4.27	MCFUCap	60
7.7.4.28	MCFUCaps	60
7.7.4.29	OpenArc	60
7.7.4.30	PreProcess	60
7.7.4.31	SetPar	60
7.7.4.32	SolveMCF	61
7.7.4.33	WriteMCF	61
8	File Documentation	63
8.1	Main.C File Reference	63
8.1.1	Detailed Description	63
8.2	MCFClass.h File Reference	65
8.2.1	Detailed Description	65
8.3	OPTUtils.h File Reference	66
8.3.1	Detailed Description	66
8.4	RelaxIV.h File Reference	67
8.4.1	Detailed Description	67

Chapter 2

Module Index

2.1 Modules

Here is a list of all modules:

Compile-time switches in MCFCClass.h	9
Classes in MCFCClass.h	10
Compile-time switches in OPTUtils.h	11
Classes in OPTUtils.h	13
Functions in OPTUtils.h	14
Compile-time switches in RelaxIV.h	15
Classes in RelaxIV.h	17

Chapter 3

Class Index

3.1 Class Hierarchy

This inheritance list is sorted roughly, but not completely, alphabetically:

MCFCClass::Eps< T >	19
MCFCClass::Inf< T >	20
MCFCClass	21
RelaxIV	52
MCFCClass::MCFException	49
MCFCClass::MCFState	50
OPTrand	51

Chapter 4

Class Index

4.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

MCFCClass::Eps< T >	19
MCFCClass::Inf< T >	20
MCFCClass	21
MCFCClass::MCFException	49
MCFCClass::MCFState	50
OPTrand	51
RelaxIV	52

Chapter 5

File Index

5.1 File List

Here is a list of all documented files with brief descriptions:

Main.C	63
MCFCClass.h	65
OPTUtils.h	66
RelaxIV.h	67

Chapter 6

Module Documentation

6.1 Compile-time switches in MCFClass.h

Defines

- `#define USERNAME0 0`

6.1.1 Detailed Description

These macros control some important details of the class interface. Although using macros for activating features of the interface is not very C++, switching off some unused features may allow some implementation to be more efficient in running time or memory.

6.1.2 Define Documentation

6.1.2.1 `#define USERNAME0 0`

Decides if 0 or 1 is the "name" of the first node. If `USERNAME0 == 1`, (warning: it has to be *exactly* 1), then the node names go from 0 to $n - 1$, otherwise from 1 to n . Note that this does not affect the position of the deficit in the deficit vectors, i.e., the deficit of the i -th node - be its "name" 'i' or 'i - 1' - is always in the i -th position of the vector.

6.2 Classes in MCFClass.h

Classes

- class [MCFClass](#)

6.3 Compile-time switches in OPTUtils.h

Defines

- #define `OPT_USE_NAMESPACES` 0
- #define `OPT_TIMERS` 5
- #define `OPT_RANDOM` 1

6.3.1 Detailed Description

These macros control how the classes OPTTimers and `OPTrand` are implemented; choose the appropriate value for your environment, or program a new version if no value suits you. Also, namespaces can be eliminated if they create problems.

6.3.2 Define Documentation

6.3.2.1 #define OPT_RANDOM 1

The class `OPTrand` is defined below to give an abstract interface to the different random generators that are used in different platforms. This is needed since random generators are one of the less standard parts of the C++ library. The value of the `OPT_RANDOM` constant selects among the different timing routines:

- 0 = an hand-made implementation of a rather good random number generator is used; note that this assumes that long ints ≥ 32 bits
- 1 = standard `rand()` / `srand()` pair, common to all C libraries but not very sophisticated
- 2 = `drand48()` / `srand48()`, common on Unix architectures and pretty good.

Any unsupported value would simply make the functions to report constant zero, which is not nice but useful to quickly fix problems if you don't use random numbers at all.

6.3.2.2 #define OPT_TIMERS 5

The class OPTtimers is defined below to give an abstract interface to the different timing routines that are used in different platforms. This is needed since time-related functions are one of the less standard parts of the C++ library. The value of the `OPT_TIMERS` constant selects among the different timing routines:

- 1 = Use the Unix `times()` routine in `sys/times.h`
- 2 = As 1 but uses `sys/timeb.h` (typical for Microsoft(TM) compilers)
- 3 = Still use `times()` of `sys/times.h`, but return wallclock time rather than CPU time
- 4 = As 3 but uses `sys/timeb.h` (typical for Microsoft(TM) compilers)
- 5 = return the user time obtained with ANSI C `clock()` function; this may result in more accurate running times w.r.t. but may be limited to ~ 72 hours on systems where ints are 32bits.

- 6 = Use the Unix `gettimeofday()` routine of `sys/time.h`.

Any unsupported value would simply make the class to report constant zero as the time.

The values 1 .. 4 rely on the constant `CLK_TCK` for converting between clock ticks and seconds; for the case where the constant is not defined by the compiler -- should not happen, but it does -- or it is defined in a wrong way, the constant is re-defined below.

6.3.2.3 `#define OPT_USE_NAMESPACES 0`

Setting `OPT_USE_NAMESPACES == 0` should instruct all codes that use `OPTUtils` stuff to avoid using namespaces; to start with, the common namespace `OPTUtils_di_unipi_it`, that contains all the types defined herein, is **not** defined.

6.4 Classes in OPTUtils.h

Classes

- class [OPTrand](#)

6.5 Functions in OPTUtils.h

Functions

- `template<class T >`
`void DfltDfInpt (istream *iStrm, T &Param, const T Dflt, const char cmntc= '#')`

6.5.1 Function Documentation

6.5.1.1 `template<class T > void DfltDfInpt (istream * iStrm, T & Param, const T Dflt, const char cmntc = '#') [inline]`

Template function for reading parameters from a istream. The function is "safe" because it works also if the istream is not given, is not be long enough or contains erroneous things.

Given a &istream (possibly NULL), `DfltDfInpt()` attempts to read Param out of it, skipping any line that begins with the comment carachter (defaulted to '#'), any blank line and any line starting with anything that can not be interpreted as a 'T'. If, for any reason, the read operation fails, then the parameter is given the default value 'Dflt'. Otherwise, all the rest of the line up to the nearest newline ('

') carachter is flushed.

Note

lines should not be longer than 255 carachters.

6.6 Compile-time switches in RelaxIV.h

Defines

- `#define SAME_GRPH_RIV 0`
- `#define DYNMC_MCF_RIV 0`
- `#define AUCTION 0`
- `#define RELAXIV_STATISTICS 0`

6.6.1 Detailed Description

These macros control some important details of the implementation. Although using macros for activating features of the implementation is not very C++, switching off some unused features may make the code more efficient in running time or memory.

6.6.2 Define Documentation

6.6.2.1 `#define AUCTION 0`

Decides if the auction/shortest paths initialization procedure is used. If `AUCTION == 1`, then an auction/shortest paths initialization procedure is provided [see `SetAuction()` below] that has been reported to make the [RelaxIV](#) algorithm run faster on some classes of instances. The auction initialization essentially "spreads" the imbalances around the graph by performing some steps of the "pure" epsilon-relaxation method: this should produce "short" augmenting steps, that seem to be the best situation for [RelaxIV](#).

By setting `AUCTION == 0`, some memory is saved.

6.6.2.2 `#define DYNMC_MCF_RIV 0`

Decides if the graph topology (arcs, nodes) can be changed. If `DYNMC_MCF_RIV > 0`, the methods of the public interface of class that allow to change the topology of the underlying network are actually implemented. Possible values of this macro are:

- `0 =>` the topology of the graph cannot be changed;
- `1 =>` the methods that "close" arcs and delete nodes are implemented;
- `2 =>` the methods that "open" previously closed arcs and add nodes are implemented;
- `3 =>` the methods that change the start and end node of a (possibly "closed") arc, delete and create new arcs are implemented.

As long as `DYNMC_MCF_RIV < 3`, `SAME_GRPH_RIV` can possibly be `> 0`; less and less data structures can be shared as `DYNMC_MCF_RIV` increases. `DYNMC_MCF_RIV == 3` implies `SAME_GRPH_RIV == 0`.

6.6.2.3 `#define RELAXIV_STATISTICS 0`

If `RELAXIV_STATISTICS > 0`, then statistic information about the behaviour of the Relaxation algorithm is computed.

6.6.2.4 `#define SAME_GRP_H_RIV 0`

Decides if all [MCFClass](#) instances share the same graph. If `SAME_GRP_H_RIV > 0`, then all the instances of the class will work on the same "topological" network, while the costs, capacities and supplies can change from one instance to another. This allows implementations to share some data structures describing the graph, e.g. by declaring them "static", saving memory when multiple instances of the solver are active at the same time. However, this also obviously inhibits some type of changes in the topology of the graph [see `DYNMC_MCF_RIV` below].

6.7 Classes in RelaxIV.h

Classes

- class [RelaxIV](#)

Chapter 7

Class Documentation

7.1 MCFClass::Eps< T > Class Template Reference

```
#include <MCFClass.h>
```

Public Member Functions

- `operator T ()`

7.1.1 Detailed Description

template<typename T> class MCFClass::Eps< T >

Very small class to simplify extracting the "machine epsilon" for a basic type (FNumber, CNumber); just use `Eps<type>()`.

The documentation for this class was generated from the following file:

- [MCFClass.h](#)

7.2 MCFClass::Inf< T > Class Template Reference

```
#include <MCFClass.h>
```

Public Member Functions

- `operator T ()`

7.2.1 Detailed Description

template<typename T> class MCFClass::Inf< T >

Very small class to simplify extracting the "+ infinity" value for a basic type (FNumber, CNumber, Index); just use `Inf<type>()`.


The documentation for this class was generated from the following file:

- [MCFClass.h](#)

7.3 MCFClass Class Reference

```
#include <MCFClass.h>
```

Inheritance diagram for MCFClass:



Classes

- class [Eps](#)
- class [Inf](#)
- class [MCFException](#)
- class [MCFState](#)

Public Types

Public types

The [MCFClass](#) defines four main public types:

- *Index*, the type of arc and node indices;
- *FNumber*, the type of flow variables, arc capacities, and node deficits;
- *CNumber*, the type of flow costs, node potentials, and arc reduced costs;
- *FONumber*, the type of objective function value.

By re-defining the types in this section, most *MCFSolver* should be made to work with any reasonable choice of data type (= one that is capable of properly representing the data of the instances to be solved). This may be relevant due to an important property of MCF problems: **if all arc capacities and node deficits are integer, then there exists an integral optimal primal solution**, and **if all arc costs are integer, then there exists an integral optimal dual solution**. Even more importantly, **many solution algorithms will in fact produce an integral primal/dual solution for free**, because **every primal/dual solution they generate during the solution process is naturally integral**. Therefore, one can use integer data types to represent everything connected with flows and/or costs if the corresponding data is integer in all instances one needs to solve. This directly translates in significant memory savings and/or speed improvements.

It is the user's responsibility to ensure that these types are set to reasonable values*. So, the experienced user may want to experiment with setting this data properly if memory footprint and/or speed is a primary concern. Note, however, that **not all solution algorithms will happily accept integer data**; one example are Interior-Point approaches, which require both flow and cost variables to be continuous (float). So, the viability of setting integer data (as well as its impact on performances) is strictly related to the specific kind of algorithm used. Since these types are common to all derived classes, they have to be set taking into account the needs of all the solvers that are going to be used, and adapting to the "worst case"; of course, *FNumber == CNumber == double* is going to always be an acceptable "worst case" setting. [MCFClass](#) may in a future be defined as a template class, with these as template parameters, but this is currently deemed overkill and avoided.

Finally, note that the above integrality property only holds for *linear* MCF problems. If any arc has a nonzero quadratic cost coefficient, optimal flows and potentials may be fractional even if all the data of the problem (comprised quadratic cost coefficients) is integer. Hence, for *quadratic* MCF solvers, a setting like `FNumber == CNumber == double` is actually mandatory*, for any reasonable algorithm will typically misbehave otherwise.

- enum `MCFParam` {
`kMaxTime` = 0, `kMaxIter`, `kEpsFlw`, `kEpsDfct`,
`kEpsCst`, `kReopt`, `kLastParam` }
- enum `MCFStatus` {
`kUnSolved` = -1, `kOK` = 0, `kStopped`, `kUnfeasible`,
`kUnbounded`, `kError` }
- enum `MCFAnswer` { `kNo` = 0, `kYes` }
- enum `MCFIFrmt` { `kDimacs` = 0, `kQDimacs`, `kMPS`, `kFWMPS` }
- typedef unsigned int `Index`
index of a node or arc (>= 0)
- typedef `Index * Index_Set`
set (array) of indices
- typedef const `Index cIndex`
a read-only index
- typedef `cIndex * cIndex_Set`
read-only index array
- typedef int `SIndex`
index of a node or arc
- typedef `SIndex * SIndex_Set`
set (array) of indices
- typedef const `SIndex cSIndex`
a read-only index
- typedef `cSIndex * cSIndex_Set`
read-only index array
- typedef double `FNumber`
type of arc flow
- typedef `FNumber * FRow`
vector of flows
- typedef const `FNumber cFNumber`
a read-only flow
- typedef `cFNumber * cFRow`
read-only flow array
- typedef double `CNumber`
type of arc flow cost

- typedef [CNumber](#) * [CRow](#)
vector of costs
- typedef const [CNumber](#) [cCNumber](#)
a read-only cost
- typedef [cCNumber](#) * [cCRow](#)
read-only cost array
- typedef double [FONumber](#)
- typedef const [FONumber](#) [cFONumber](#)
a read-only o.f. value
- typedef [MCFState](#) * [MCFStatePtr](#)
pointer to a [MCFState](#)

Public Member Functions

Constructors

- [MCFClass](#) ([cIndex](#) nmX=0, [cIndex](#) mmX=0)

Other initializations

- virtual void [LoadNet](#) ([cIndex](#) nmX=0, [cIndex](#) mmX=0, [cIndex](#) pn=0, [cIndex](#) pm=0, [cFRow](#) pU=NULL, [cCRow](#) pC=NULL, [cFRow](#) pDfct=NULL, [cIndex_Set](#) pSn=NULL, [cIndex_Set](#) pEn=NULL)=0
- virtual void [LoadDMX](#) (istream &DMXs, bool IsQuad=false)
- virtual void [PreProcess](#) (void)
- virtual void [SetPar](#) (int par, int val)
- virtual void [SetPar](#) (int par, double val)
- virtual void [GetPar](#) (int par, int &val)
- virtual void [GetPar](#) (int par, double &val)
- virtual void [SetMCFTime](#) (bool TimeIt=true)

Solving the problem

- virtual void [SolveMCF](#) (void)=0
- int [MCFGetStatus](#) (void)

Reading flow solution

- virtual void [MCFGetX](#) ([FRow](#) F, [Index_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=[Inf](#)< [Index](#) >())=0
- virtual [cFRow](#) [MCFGetX](#) (void)
- virtual bool [HaveNewX](#) (void)

Reading potentials

- virtual void [MCFGetPi](#) ([CRow](#) P, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=[Inf](#)< [Index](#) >())=0
- virtual [cCRow](#) [MCFGetPi](#) (void)
- virtual bool [HaveNewPi](#) (void)

Reading reduced costs

- virtual void [MCFGetRC](#) ([CRow](#) CR, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf<[Index](#) >())=0
- virtual [cCRow](#) [MCFGetRC](#) (void)
- virtual [CNumber](#) [MCFGetRC](#) ([cIndex](#) i)=0

Reading the objective function value

- virtual [FONumber](#) [MCFGetFO](#) (void)=0
- virtual [FONumber](#) [MCFGetDFO](#) (void)

Getting unfeasibility certificate

- virtual [FNumber](#) [MCFGetUnfCut](#) ([Index_Set](#) Cut)

Getting unboundedness certificate

- virtual [Index](#) [MCFGetUnbCycl](#) ([Index_Set](#) Pred, [Index_Set](#) ArcPred)

Saving/restoring the state of the solver

- virtual [MCFStatePtr](#) [MCFGetState](#) (void)
- virtual void [MCFPutState](#) ([MCFStatePtr](#) S)

Time the code

- void [TimeMCF](#) (double &t_us, double &t_ss)
- double [TimeMCF](#) (void)

Check the solutions

- void [CheckPSol](#) (void)
- void [CheckDSol](#) (void)

Reading graph size

- [Index](#) [MCFnmax](#) (void)
- [Index](#) [MCFmmax](#) (void)
- [Index](#) [MCFn](#) (void)
- [Index](#) [MCFm](#) (void)

Reading graph topology

- virtual void [MCFArcs](#) ([Index_Set](#) Startv, [Index_Set](#) Endv, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf<[Index](#) >())=0
- virtual [Index](#) [MCFSNde](#) ([cIndex](#) i)=0
- virtual [Index](#) [MCFENde](#) ([cIndex](#) i)=0
- virtual [cIndex_Set](#) [MCFSNdes](#) (void)
- virtual [cIndex_Set](#) [MCFENdes](#) (void)

Reading arc costs

- virtual void [MCFCosts](#) ([CRow](#) Costv, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf<[Index](#) >())=0

- virtual [CNumber MCFCost](#) ([cIndex](#) i)=0
- virtual [cCRow MFCFCosts](#) (void)
- virtual void [MCFQCoef](#) ([CRow](#) Qv, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- virtual [CNumber MCFQCoef](#) ([cIndex](#) i)
- virtual [cCRow MCFQCoef](#) (void)

Reading arc capacities

- virtual void [MCFUCaps](#) ([FRow](#) UCapv, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual [FNumber MCFUCap](#) ([cIndex](#) i)=0
- virtual [cFRow MCFUCaps](#) (void)

Reading node deficits

- virtual void [MCFDfcts](#) ([FRow](#) Dfctv, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual [FNumber MCFDfct](#) ([cIndex](#) i)=0
- virtual [cFRow MCFDfcts](#) (void)

Write problem to file

- virtual void [WriteMCF](#) (ostream &oStrm, int frmt=0)

Changing the costs

- virtual void [ChgCosts](#) ([cCRow](#) NCost, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual void [ChgCost](#) ([Index](#) arc, [CNumber](#) NCost)=0
- virtual void [ChgQCoef](#) ([cCRow](#) NQCoef=NULL, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())
- virtual void [ChgQCoef](#) ([Index](#) arc, [CNumber](#) NQCoef)

Changing the capacities

- virtual void [ChgUCaps](#) ([cFRow](#) NCap, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual void [ChgUCap](#) ([Index](#) arc, [FNumber](#) NCap)=0

Changing the deficits

- virtual void [ChgDfcts](#) ([cFRow](#) NDfct, [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=Inf< [Index](#) >())=0
- virtual void [ChgDfct](#) ([Index](#) node, [FNumber](#) NDfct)=0

Changing graph topology

- virtual void [CloseArc](#) ([cIndex](#) name)=0
- virtual bool [IsClosedArc](#) ([cIndex](#) name)=0
- virtual void [DelNode](#) ([cIndex](#) name)=0
- virtual void [OpenArc](#) ([cIndex](#) name)=0
- virtual [Index AddNode](#) ([FNumber](#) aDfct)=0
- virtual void [ChangeArc](#) ([cIndex](#) name, [cIndex](#) nSN=Inf< [Index](#) >(), [cIndex](#) nEN=Inf< [Index](#) >())=0

- virtual void [DelArc](#) ([cIndex](#) name)=0
- virtual bool [IsDeletedArc](#) ([cIndex](#) name)=0
- virtual [Index AddArc](#) ([cIndex](#) Start, [cIndex](#) End, [cFNumber](#) aU, [cCNumber](#) aC)=0

Destructor

- virtual [~MCFClass](#) ()

Protected Member Functions

Managing comparisons.

The following methods are provided for making it easier to perform comparisons, with and without tolerances.

- template<class T >
bool [ETZ](#) (T x, const T eps)
- template<class T >
bool [GTZ](#) (T x, const T eps)
- template<class T >
bool [GEZ](#) (T x, const T eps)
- template<class T >
bool [LTZ](#) (T x, const T eps)
- template<class T >
bool [LEZ](#) (T x, const T eps)
- template<class T >
bool [GT](#) (T x, T y, const T eps)
- template<class T >
bool [LT](#) (T x, T y, const T eps)

Protected Attributes

- [Index n](#)
total number of nodes
- [Index nmax](#)
maximum number of nodes
- [Index m](#)
total number of arcs
- [Index mmax](#)
maximum number of arcs
- int [status](#)
- bool [Sensv](#)
true <=> the latest optimal solution should be exploited
- OPTtimers * [MCFt](#)
timer for performances evaluation
- [FNumber EpsFlw](#)

precision for comparing arc flows / capacities

- [FNumber EpsDfct](#)

precision for comparing node deficits

- [CNumber EpsCst](#)

precision for comparing arc costs

- double [MaxTime](#)
- int [MaxIter](#)

7.3.1 Detailed Description

This abstract base class defines a standard interface for (linear or convex quadratic separable) Min Cost Flow (MCF) problem solvers.

The data of the problem consist of a (directed) graph $G = (N, A)$ with $n = |N|$ nodes and $m = |A|$ (directed) arcs. Each node 'i' has a deficit $b[i]$, i.e., the amount of flow that is produced/consumed by the node: source nodes (which produce flow) have negative deficits and sink nodes (which consume flow) have positive deficits. Each arc '(i, j)' has an upper capacity $U[i, j]$, a linear cost coefficient $C[i, j]$ and a (non negative) quadratic cost coefficient $Q[i, j]$. Flow variables $X[i, j]$ represents the amount of flow to be sent on arc (i, j). Parallel arcs, i.e., multiple copies of the same arc '(i, j)' (with possibly different costs and/or capacities) are in general allowed. The formulation of the problem is therefore:

$$\begin{aligned} \min \quad & \sum_{(i,j) \in A} C[i, j]X[i, j] + Q[i, j]X[i, j]^2/2 \\ (1) \quad & \sum_{(j,i) \in A} X[j, i] - \sum_{(i,j) \in A} X[i, j] = b[i] \quad i \in N \\ (2) \quad & 0 \leq X[i, j] \leq U[i, j] \quad (i, j) \in A \end{aligned}$$

The n equations (1) are the flow conservation constraints and the $2m$ inequalities (2) are the flow non-negativity and capacity constraints. At least one of the flow conservation constraints is redundant, as the demands must be balanced ($\sum_{i \in N} b[i] = 0$); indeed, exactly $n - \text{ConnectedComponents}(G)$ flow conservation constraints are redundant, as demands must be balanced in each connected component of G . Let us denote by QA and LA the disjoint subsets of A containing, respectively, "quadratic" arcs (with $Q[i, j] > 0$) and "linear" arcs (with $Q[i, j] = 0$); the (MCF) problem is linear if QA is empty, and nonlinear (convex quadratic) if QA is nonempty.

The dual of the problem is:

$$\begin{aligned} \max \quad & \sum_{i \in N} P[i]b[i] - \sum_{(i,j) \in A} W[i, j]U[i, j] - \sum_{(i,j) \in AQ} V[i, j]^2/(2 * Q[i, j]) \\ (3.a) \quad & C[i, j] - P[i] + P[j] + W[i, j] - Z[i, j] = 0 \quad (i, j) \in AL \\ (3.b) \quad & C[i, j] - P[i] + P[j] + W[i, j] - Z[i, j] = V[i, j] \quad (i, j) \in AQ \\ (4.a) \quad & W[i, j] \geq 0 \quad (i, j) \in A \\ (4.b) \quad & Z[i, j] \geq 0 \quad (i, j) \in A \end{aligned}$$

$Pi[]$ is said the vector of node potentials for the problem, $W[]$ are bound variables and $Z[]$ are slack variables. Given $Pi[]$, the quantities

$$RC[i, j] = C[i, j] + Q[i, j] * X[i, j] - P[i] + P[j]$$

are said the "reduced costs" of arcs.

A primal and dual feasible solution pair is optimal if and only if the complementary slackness conditions

$$RC[i, j] > 0 \Rightarrow X[i, j] = 0$$

$$RC[i, j] < 0 \Rightarrow X[i, j] = U[i, j]$$

are satisfied for all arcs (i, j) of A.

The [MCFClass](#) class provides an interface with methods for managing and solving problems of this kind. Actually, the class can also be used as an interface for more general NonLinear MCF problems, where the cost function either nonseparable ($C(X)$) or arc-separable ($\sum_{(i,j) \in A} C_{i,j}(X[i, j])$). However, solvers for NonLinear MCF problems are typically objective-function-specific, and there is no standard way for inputting a nonlinear function different from a separable convex quadratic one, so only the simplest form is dealt with in the interface, leaving more complex NonLinear parts to the interface of derived classes.

7.3.2 Member Typedef Documentation

7.3.2.1 typedef double MCFClass::FONumber

type of the objective function: has to hold sums of products of FNumber(s) by CNumber(s)

7.3.3 Member Enumeration Documentation

7.3.3.1 enum MCFClass::MCFAnswer

Public enum describing the possible reoptimization status of the MCF solver.

Enumerator:

kNo no
kYes yes

7.3.3.2 enum MCFClass::MCFFIFrmt

Public enum describing the possible file formats in [WriteMCF\(\)](#).

Enumerator:

kDimacs DIMACS file format for MCF.
kQDimacs quadratic DIMACS file format for MCF
kMPS MPS file format for LP.
kFWMPS "Fixed Width" MPS format

7.3.3.3 enum MCFClass::MCFParam

Public enum describing the possible parameters of the MCF solver, to be used with the methods [SetPar\(\)](#) and [GetPar\(\)](#).

Enumerator:

kMaxTime max time

kMaxIter max number of iteration

kEpsFlw tolerance for flows

kEpsDfct tolerance for deficits

kEpsCst tolerance for costs

kReopt whether or not to reoptimize

kLastParam dummy parameter: this is used to allow derived classes to "extend" the set of parameters.

7.3.3.4 enum MCFClass::MCFStatus

Public enum describing the possible status of the MCF solver.

Enumerator:

kUnsolved no solution available

kOK optimal solution found

kStopped optimization stopped

kUnfeasible problem is unfeasible

kUnbounded problem is unbounded

kError error in the solver

7.3.4 Constructor & Destructor Documentation

7.3.4.1 MCFClass::MCFClass (cIndex *nmx* = 0, cIndex *mmx* = 0) [inline]

Constructor of the class.

nmx and *mmx*, if provided, are taken to be respectively the maximum number of nodes and arcs in the network. If nonzero values are passed, memory allocation can be anticipated in the constructor, which is sometimes desirable. The maximum values are stored in the protected fields *nmax* and *mmax*, and can be changed with [LoadNet\(\)](#) [see below]; however, changing them typically requires memory allocation/deallocation, which is sometimes undesirable outside the constructor.

After that an object has been constructed, no problem is loaded; this has to be done with [LoadNet\(\)](#) [see below]. Thus, it is an error to invoke any method which requires the presence of a problem (typicall all except those in the initializations part). The base class provides two protected fields *n* and *m* for the current number of nodes and arcs, respectively, that are set to 0 in the constructor precisely to indicate that no instance is currently loaded.

7.3.4.2 virtual MCFClass::~~MCFClass () [inline, virtual]

Destructor of the class. The implementation in the base class only deletes the MCFt field. It is virtual, as it should be.

7.3.5 Member Function Documentation

7.3.5.1 **virtual Index MCFClass::AddArc (cIndex *Start*, cIndex *End*, cFNumber *aU*, cCNumber *aC*) [pure virtual]**

Add the new arc (*Start* , *End*) with cost *aC* and capacity *aU*, returning its name. `Inf<Index>()` is returned if there is no room for a new arc. Remember that arc names go from 0 to `mmax - 1`.

Implemented in [RelaxIV](#).

7.3.5.2 **virtual Index MCFClass::AddNode (cFNumber *aDfct*) [pure virtual]**

Add a new node with deficit *aDfct*, returning its name. `Inf<Index>()` is returned if there is no room for a new node. Remember that the node names are either { 0 .. `nmax - 1` } or { 1 .. `nmax` }, depending on the value of `USNAME0`.

Implemented in [RelaxIV](#).

7.3.5.3 **virtual void MCFClass::ChangeArc (cIndex *name*, cIndex *nSN* = `Inf< Index >()`, cIndex *nEN* = `Inf< Index >()`) [pure virtual]**

Change the starting and/or ending node of arc '*name*' to *nSN* and *nEN*. Each parameter being `Inf<Index>()` means to leave the previous starting or ending node untouched. When this method is called '*name*' can be either the name of a "normal" arc or that of a "closed" arc [see [CloseArc\(\)](#) above]: in the latter case, at the end of [ChangeArc\(\)](#) the arc is **still closed**, and it remains so until `OpenArc(name)` [see above] is called.

Implemented in [RelaxIV](#).

7.3.5.4 **void MCFClass::CheckDSol (void) [inline]**

Check that the dual solution returned by the solver is dual feasible. (to within the tolerances set by `SetPar(kEps****)` [see above], if any). Also, check that the objective function value is correct.

This method is implemented by the base class, using the above methods for collecting the solutions and the methods of the next section for reading the data of the problem; as such, they will work for any derived class that properly implements all these methods.

7.3.5.5 **void MCFClass::CheckPSol (void) [inline]**

Check that the primal solution returned by the solver is primal feasible. (to within the tolerances set by `SetPar(kEps****)` [see above], if any). Also, check that the objective function value is correct.

This method is implemented by the base class, using the above methods for collecting the solutions and the methods of the next section for reading the data of the problem; as such, they will work for any derived class that properly implements all these methods.

7.3.5.6 **virtual void MCFClass::ChgCost (Index *arc*, cCNumber *NCost*) [pure virtual]**

Change the cost of the *i*-th arc.

Note

changing the costs of arcs that **do not exist** is **not allowed**; only arcs which have not been

closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C_INF costs] can be touched with these methods.

7.3.5.7 virtual void MCFClass::ChgCosts (cCRow *NCost*, cIndex_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]

Change the arc costs. In particular, change the costs that are:

- listed in into the vector of indices ‘nms’ (ordered in increasing sense and Inf<Index>()-terminated),
- *and* whose name belongs to the interval [‘strt’, ‘stp’).

That is, if $strt \leq nms[i] < stp$, then the $nms[i]$ -th cost will be changed to $NCost[i]$. If $nms == NULL$ (as the default), *all* the entries in the given range will be changed; if $stp > MCFm()$, then the smaller bound is used.

Note

changing the costs of arcs that *do not exist* is *not allowed*; only arcs which have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C_INF costs] can be touched with these methods.

7.3.5.8 virtual void MCFClass::ChgDfct (Index *node*, cFNumber *NDfct*) [pure virtual]

Change the deficit of the i -th node.

Note that, in [ChgDfct\[s\]\(\)](#), node "names" (i , $strt$ / stp or those contained in $nms[]$) go from 0 to $n - 1$, regardless to the value of `USENAME0`; hence, if `USENAME0 == 0` then the first node is "named 1", but its deficit can be changed e.g. with [ChgDfcts](#)(0 , new_deficit).

Note

changing the capacities of nodes that *do not exist* is *not allowed*; only nodes that have not been deleted [see [DelNode\(\)](#) below] can be touched with these methods.

7.3.5.9 virtual void MCFClass::ChgDfcts (cFRow *NDfct*, cIndex_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]

Change the node deficits. In particular, change the deficits that are:

- listed in into the vector of indices ‘nms’ (ordered in increasing sense and Inf<Index>()-terminated),
- *and* whose name belongs to the interval [‘strt’, ‘stp’).

That is, if $strt \leq nms[i] < stp$, then the $nms[i]$ -th deficit will be changed to $NDfct[i]$. If $nms == NULL$ (as the default), *all* the entries in the given range will be changed; if $stp > MCFn()$, then the smaller bound is used.

Note that, in [ChgDfcts\(\)](#), node "names" ($strt$, stp or those contained in $nms[]$) go from 0 to $n - 1$, regardless to the value of `USENAME0`; hence, if `USENAME0 == 0` then the first node is "named 1", but its deficit can be changed e.g. with [ChgDfcts](#)(&new_deficit , NULL , 0 , 1).

Note

changing the capacities of nodes that **do not exist** is **not allowed**; only nodes that have not been deleted [see [DelNode\(\)](#) below] can be touched with these methods.

7.3.5.10 **virtual void MCFClass::ChgQCoef (Index *arc*, cCNumber *NQCoef*) [inline, virtual]**

Parameters

NQCoef Change the quadratic coefficient of the cost of the *i*-th arc.

Note that the method is **not** pure virtual: an implementation is provided for "pure linear" MCF solvers that only work with all zero quadratic coefficients.

Note

changing the costs of arcs that **do not exist** is **not allowed**; only arcs which have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C_INF costs] can be touched with these methods.

7.3.5.11 **virtual void MCFClass::ChgQCoef (cCRow *NQCoef* = NULL, cIndex_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf<Index>()) [inline, virtual]**

Parameters

stp Change the quadratic coefficients of the arc costs. In particular, change the coefficients that are:

- listed in into the vector of indices '*nms*' (ordered in increasing sense and Inf<Index>()-terminated),
- **and** whose name belongs to the interval [*strt*, '*stp*').

That is, if *strt* ≤ *nms*[*i*] < *stp*, then the *nms*[*i*]-th cost will be changed to *NCost*[*i*]. If *nms* == NULL (as the default), **all** the entries in the given range will be changed; if *stp* > [MCFm\(\)](#), then the smaller bound is used. If *NQCoef* == NULL, all the specified coefficients are set to zero.

Note that the method is **not** pure virtual: an implementation is provided for "pure linear" MCF solvers that only work with all zero quadratic coefficients.

Note

changing the costs of arcs that **do not exist** is **not allowed**; only arcs which have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C_INF costs] can be touched with these methods.

7.3.5.12 **virtual void MCFClass::ChgUCap (Index *arc*, cFNumber *NCap*) [pure virtual]**

Change the capacity of the *i*-th arc.

Note

changing the capacities of arcs that **do not exist** is **not allowed**; only arcs that have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C_INF costs] can be touched with these methods.

7.3.5.13 virtual void MCFClass::ChgUCaps (cFRow *NCap*, cIndex_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]

Change the arc capacities. In particular, change the capacities that are:

- listed in into the vector of indices ‘nms’ (ordered in increasing sense and Inf<Index>()-terminated),
- *and* whose name belongs to the interval [‘strt’, ‘stp’).

That is, if $strt \leq nms[i] < stp$, then the $nms[i]$ -th capacity will be changed to $NCap[i]$. If $nms == NULL$ (as the default), *all* the entries in the given range will be changed; if $stp > MCFm()$, then the smaller bound is used.

Note

changing the capacities of arcs that *do not exist* is *not allowed*; only arcs that have not been closed/deleted [see [CloseArc\(\)](#) / [DelArc\(\)](#) below and [LoadNet\(\)](#) above about C_INF costs] can be touched with these methods.

7.3.5.14 virtual void MCFClass::CloseArc (cIndex *name*) [pure virtual]

"Close" the arc ‘name’. Although all the associated information (name, cost, capacity, end and start node) is kept, the arc is removed from the problem until [OpenArc\(i\)](#) [see below] is called.

"closed" arcs always have 0 flow, but are otherwise counted as any other arc; for instance, [MCFm\(\)](#) does *not* decrease as an effect of a call to [CloseArc\(\)](#). How this closure is implemented is solver-specific.

Implemented in [RelaxIV](#).

7.3.5.15 virtual void MCFClass::DelArc (cIndex *name*) [pure virtual]

Delete the arc ‘name’. Unlike "closed" arcs, all the information associated with a deleted arc is lost and ‘name’ is made available as a name for new arcs to be created with [AddArc\(\)](#) [see below].

It furthermore ‘name’ is the last arc, the number of arcs as reported by [MCFm\(\)](#) is reduced by at least one, until the m-th arc is not a deleted one. Otherwise, the flow on the arc is always ensured to be 0.

Implemented in [RelaxIV](#).

7.3.5.16 virtual void MCFClass::DelNode (cIndex *name*) [pure virtual]

Delete the node ‘name’.

For any value of ‘name’, all incident arcs to that node are closed [see [CloseArc\(\)](#) above] (*not* Deleted, see [DelArc\(\)](#) below) and the deficit is set to zero.

It furthermore ‘name’ is the last node, the number of nodes as reported by [MCFn\(\)](#) is reduced by at least one, until the n-th node is not a deleted one.

Implemented in [RelaxIV](#).

7.3.5.17 template<class T> bool MCFClass::ETZ (T *x*, const T *eps*) [inline, protected]

true if flow *x* is equal to zero (possibly considering tolerances).

7.3.5.18 void MCFClass::GetPar (int *par*, double & *val*) [inline, virtual]

This method returns one of the integer parameter of the algorithm.

Parameters

par is the parameter to return [see SetPar(double) for comments];

val upon return, it will contain the value of the parameter.

The base class implementation handles the parameters kEpsFlw, kEpsDfct, kEpsCst, and kMaxTime.

7.3.5.19 void MCFClass::GetPar (int *par*, int & *val*) [inline, virtual]

This method returns one of the integer parameter of the algorithm.

Parameters

par is the parameter to return [see SetPar(int) for comments];

val upon return, it will contain the value of the parameter.

The base class implementation handles the parameters kMaxIter and kReopt.

Reimplemented in [RelaxIV](#).

7.3.5.20 template<class T > bool MCFClass::GEZ (T *x*, const T *eps*) [inline, protected]

true if flow *x* is greater than or equal to zero (possibly considering tolerances).

7.3.5.21 template<class T > bool MCFClass::GT (T *x*, T *y*, const T *eps*) [inline, protected]

true if flow *x* is greater than flow *y* (possibly considering tolerances).

7.3.5.22 template<class T > bool MCFClass::GTZ (T *x*, const T *eps*) [inline, protected]

true if flow *x* is greater than zero (possibly considering tolerances).

7.3.5.23 virtual bool MCFClass::HaveNewPi (void) [inline, virtual]

Return true if a different (approximately) optimal dual solution is available. If the method returns true, then any subsequent call to (any form of) [MCFGetPi\(\)](#) will return a different dual solution, and [MCFGetRC\(\)](#) [see below] will return the corresponding reduced costs. The new solution need not be optimal (although, ideally, it has to be "good"); this can be checked by comparing its objective function value, that will be returned by a call to [MCFGetDFO\(\)](#) [see below].

Any subsequent call of [HaveNewPi\(\)](#) that returns true produces a new solution, until the first that returns false; from then on, no new solutions will be generated until something changes in the problem's data.

Note that a default implementation of [HaveNewPi\(\)](#) is provided which is good for those solvers that only produce one optimal dual solution.

7.3.5.24 virtual bool MCFClass::HaveNewX (void) [inline, virtual]

Return true if a different (approximately) optimal primal solution is available. If the method returns true, then any subsequent call to (any form of) [MCFGetX\(\)](#) will return a different primal solution w.r.t. the one that was being returned *before* the call to [HaveNewX\(\)](#). This solution need not be optimal (although, ideally, it has to be "good"); this can be checked by comparing its objective function value, that will be returned by a call to [MCFGetFO\(\)](#) [see below].

Any subsequent call of [HaveNewX\(\)](#) that returns true produces a new solution, until the first that returns false; from then on, no new solutions will be generated until something changes in the problem's data.

Note that a default implementation of [HaveNewX\(\)](#) is provided which is good for those solvers that only produce one optimal primal solution.

7.3.5.25 virtual bool MCFClass::IsClosedArc (cIndex name) [pure virtual]

[IsClosedArc\(\)](#) returns true if and only if the arc 'name' is closed.

Implemented in [RelaxIV](#).

7.3.5.26 virtual bool MCFClass::IsDeletedArc (cIndex name) [pure virtual]

Return true if and only if the arc 'name' is deleted. It should only be called with name < [MCFm\(\)](#), as every other arc is deleted by definition.

Implemented in [RelaxIV](#).

7.3.5.27 template<class T > bool MCFClass::LEZ (T x, const T eps) [inline, protected]

true if flow x is less than or equal to zero (possibly considering tolerances).

7.3.5.28 void MCFClass::LoadDMX (istream & DMXs, bool IsQuad = false) [inline, virtual]

Read a MCF instance in DIMACS standard format from the istream. The format is the following. The first line must be

```
p min <number of="" nodes>=""> <number of="" arcs>="">
```

Then the node definition lines must be found, in the form

```
n <node number>=""> <node supply>="">
```

Not all nodes need have a node definition line; these are given zero supply, i.e., they are transshipment nodes (supplies are the inverse of deficits, i.e., a node with positive supply is a source node). Finally, the arc definition lines must be found, in the form

```
a <start node>=""> <end node>=""> <lower bound>=""> <upper bound>=""> <flow cost>="">
```

There must be exactly <number of="" arcs>=""> arc definition lines in the file.

This method is *not* pure virtual because an implementation is provided by the base class, using the [LoadNet\(\)](#) method (which *is* pure virtual). However, the method *is* virtual to allow derived classes to implement more efficient versions, should they have any reason to do so.

Note

Actually, the file format accepted by LoadDMX (at least in the base class implementation) is more general than the DIMACS standard format, in that it is allowed to mix node and arc definitions in any order, while the DIMACS file requires all node information to appear before all arc information.

Other than for the above, this method is assumed to allow for quadratic* Dimacs files, encoding for convex quadratic separable Min Cost Flow instances. This is a simple extension where each arc descriptor has a sixth field, `<quadratic cost>=""`. The provided istream is assumed to be quadratic Dimacs file if `IsQuad` is true, and a regular linear Dimacs file otherwise.

7.3.5.29 virtual void MCFClass::LoadNet (cIndex nm_x = 0, cIndex mm_x = 0, cIndex pn = 0, cIndex pm = 0, cFRow pU = NULL, cCRow pC = NULL, cFRow pDfct = NULL, cIndex_Set pSn = NULL, cIndex_Set pEn = NULL) [pure virtual]

Inputs a new network.

The parameters nm_x and mm_x are the new max number of nodes and arcs, possibly overriding those set in the constructor [see above], although at the likely cost of memory allocation and deallocation. Passing nm_x == mm_x == 0 is intended as a signal to the solver to deallocate everything and wait for new orders; in this case, all the other parameters are ignored.

Otherwise, in principle all the other parameters have to be provided. Actually, some of them may not be needed for special classes of MCF problems (e.g., costs in a MaxFlow problem, or start/end nodes in a problem defined over a graph with fixed topology, such as a complete graph). Also, passing NULL is allowed to set default values.

The meaning of the parameters is the following:

- pn is the current number of nodes of the network (\leq nmax).
- pm is the number of arcs of the network (\leq mmax).
- pU is the m-vector of the arc upper capacities; capacities must be nonnegative, but can in principle be infinite ($= F_INF$); passing pU == NULL means that all capacities are infinite;
- pC is the m-vector of the arc costs; costs must be finite ($< C_INF$); passing pC == NULL means that all costs must be 0.
- pDfct is the n-vector of the node deficits; source nodes have negative deficits and sink nodes have positive deficits; passing pDfct == NULL means that all deficits must be 0 (a circulation problem);
- pSn is the m-vector of the arc starting nodes; pSn == NULL is in principle not allowed, unless the topology of the graph is fixed;
- pEn is the m-vector of the arc ending nodes; same comments as for pSn.

Note that node "names" in the arrays pSn and pEn must go from 1 to pn if the macro USANAME0 [see above] is set to 0, while they must go from 0 to pn - 1 if USANAME0 is set to 1. In both cases, however, the deficit of the first node is read from the first (0-th) position of pDfct, that is if USANAME0 == 0 then the deficit of the node with name 'i' is read from pDfct[i - 1].

The data passed to [LoadNet\(\)](#) can be used to specify that the arc 'i' must not "exist" in the problem. This is done by passing pC[i] == C_INF; solvers which don't read costs are forced to read them in order to check this, unless they provide alternative solver-specific ways to accomplish the same tasks. These arcs

are "closed", as for the effect of [CloseArc\(\)](#) [see below]. "invalid" costs ($== C_INF$) are set to 0 in order to being subsequently capable of "opening" them back with [OpenArc\(\)](#) [see below]. The way in which these non-existent arcs are physically dealt with is solver-specific; in some solvers, for instance, this could be obtained by simply putting their capacity to zero. Details about these issues should be found in the interface of derived classes.

Note that the quadratic part of the objective function, if any, is not dealt with in [LoadNet\(\)](#); it can only be separately provided with [ChgQCoef\(\)](#) [see below]. By default, the problem is linear, i.e., all coefficients of the second-order terms in the objective function are assumed to be zero.

Implemented in [RelaxIV](#).

7.3.5.30 **template<class T> bool MCFClass::LT (T x, T y, const T eps) [inline, protected]**

true if flow x is less than flow y (possibly considering tolerances).

7.3.5.31 **template<class T> bool MCFClass::LTZ (T x, const T eps) [inline, protected]**

true if flow x is less than zero (possibly considering tolerances).

7.3.5.32 **virtual void MCFClass::MCFArcs (Index_Set Startv, Index_Set Endv, cIndex_Set nms = NULL, cIndex strt = 0, Index stp = Inf<Index>()) [pure virtual]**

Write the starting (tail) and ending (head) nodes of the arcs in Startv[] and Endv[]. If nms == NULL, then the information relative to all arcs is written into Startv[] and Endv[], otherwise Startv[i] and Endv[i] contain the information relative to arc nms[i] (nms[] must be Inf<Index>()-terminated).

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with $strt \leq i < \min(MCFm(), stp)$. 'strt' and 'stp' work in "&&" with nms; that is, if nms != NULL then only the values corresponding to arcs which are *both* in nms and whose index is in the correct range are returned.

Startv or Endv can be NULL, meaning that only the other information is required.

Note

If USERNAME0 == 0 then the returned node names will be in the range 1 .. n, while if USERNAME0 == 1 the returned node names will be in the range 0 .. n - 1.

If the graph is "dynamic", be careful to use [MCFn\(\)](#) e [MCFm\(\)](#) to properly choose the dimension of nodes and arcs arrays.

7.3.5.33 **virtual CNumber MCFClass::MCFCost (cIndex i) [pure virtual]**

Return the cost of the i-th arc.

Implemented in [RelaxIV](#).

7.3.5.34 **virtual cCRow MCFClass::MCFCosts (void) [inline, virtual]**

Return a read-only pointer to an internal vector containing the arc costs. Since this may *not* always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented in [RelaxIV](#).

7.3.5.35 virtual void MCFClass::MCFCosts (CRow Costv, cIndex_Set nms = NULL, cIndex strt = 0, Index stp = Inf< Index >()) [pure virtual]

Write the arc costs into Costv[]. If nms == NULL, then all the costs are written, otherwise Costv[i] contains the information relative to arc nms[i] (nms[] must be Inf<Index>()-terminated).

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with $strt \leq i < \min(\text{MCFm}(), \text{stp})$. 'strt' and 'stp' work in "&&" with nms; that is, if nms != NULL then only the values corresponding to arcs which are *both* in nms and whose index is in the correct range are returned.

7.3.5.36 virtual FNumber MCFClass::MCFDfct (cIndex i) [pure virtual]

Return the deficit of the i-th node.

Note

Here node "names" go from 0 to n - 1, regardless to the value of USERNAME0; hence, if USERNAME0 == 0 then the first node is "named 1", but its deficit is returned by MCFDfct(0).

Implemented in [RelaxIV](#).

7.3.5.37 virtual cFRow MCFClass::MCFDfcts (void) [inline, virtual]

Return a read-only pointer to an internal vector containing the node deficits. Since this may *not always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Note

Here node "names" go from 0 to n - 1, regardless to the value of USERNAME0; hence, if USERNAME0 == 0 then the first node is "named 1", but its deficit is contained in [MCFDfcts](#)(0).

Reimplemented in [RelaxIV](#).

7.3.5.38 virtual void MCFClass::MCFDfcts (FRow Dfctv, cIndex_Set nms = NULL, cIndex strt = 0, Index stp = Inf< Index >()) [pure virtual]

Write the node deficits into Dfctv[]. If nms == NULL, then all the deficits are written, otherwise Dfctv[i] contains the information relative to node nms[i] (nms[] must be Inf<Index>()-terminated).

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the nodes 'i' with $strt \leq i < \min(\text{MCFn}(), \text{stp})$. 'strt' and 'stp' work in "&&" with nms; that is, if nms != NULL then only the values corresponding to nodes which are *both* in nms and whose index is in the correct range are returned.

Note

Here node "names" (strt and stp, those contained in nms[] or 'i' in [MCFDfct](#)()) go from 0 to n - 1, regardless to the value of USERNAME0; hence, if USERNAME0 == 0 then the first node is "named 1", but its deficit is returned by MCFDfcts(Dfctv , NULL , 0 , 1).

7.3.5.39 virtual Index MCFClass::MCFENde (cIndex i) [pure virtual]

Return the ending (head) node of the arc 'i'.

Note

If USERNAME0 == 0 then the returned node names will be in the range 1 .. n, while if USERNAME0 == 1 the returned node names will be in the range 0 .. n - 1.

Implemented in [RelaxIV](#).

7.3.5.40 virtual cIndex_Set MCFClass::MCFENdes (void) [inline, virtual]

Return a read-only pointer to an internal vector containing the ending (head) nodes for each arc. Since this may *not always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented in [RelaxIV](#).

7.3.5.41 virtual FNumber MCFClass::MCFGetDFO (void) [inline, virtual]

Return the objective function value of the dual solution currently returned by [MCFGetPi\(\)](#) / [MCFGetRC\(\)](#). This value (possibly) changes after any call to [HaveNewPi\(\)](#) that returns true. The relations between [MCFGetStatus\(\)](#) and [MCFGetDFO\(\)](#) are analogous to these of [MCFGetFO\(\)](#), except that a finite value corresponding to kStopped must be a lower bound on the optimal objective function value (typically, the objective function value one dual feasible solution).

A default implementation is provided for [MCFGetDFO\(\)](#), which is good for MCF solvers where the primal and dual optimal solution values always are identical (except if the problem is unfeasible/unbounded).

7.3.5.42 virtual FNumber MCFClass::MCFGetFO (void) [pure virtual]

Return the objective function value of the primal solution currently returned by [MCFGetX\(\)](#).

If [MCFGetStatus\(\)](#) == kOK, this is guaranteed to be the optimal objective function value of the problem (to within the optimality tolerances), but only prior to any call to [HaveNewX\(\)](#) that returns true. [MCFGetFO\(\)](#) typically returns Inf<FNumber>() if [MCFGetStatus\(\)](#) == kUnfeasible and

- Inf<FNumber>() if [MCFGetStatus\(\)](#) == kUnbounded. If [MCFGetStatus\(\)](#) == kStopped and [MCFGetFO\(\)](#) returns a finite value, it must be an upper bound on the optimal objective function value (typically, the objective function value of one primal feasible solution).

Implemented in [RelaxIV](#).

7.3.5.43 virtual cCRow MCFClass::MCFGetPi (void) [inline, virtual]

Return a read-only pointer to an internal data structure containing the node potentials. Since this may *not always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented in [RelaxIV](#).

7.3.5.44 **virtual void MCFClass::MCFGetPi (CRow *P*, cIndex_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]**

Writes the optimal node potentials in the vector *P*[]. If *nms* == NULL, the node potential of node 'i' (i in 0 .. *n* - 1) is written in *P*[*i*] (note that here node names always start from zero, regardless to the value of *USENAME0*). If *nms* != NULL, it must point to a vector of indices in 0 .. *n* - 1 (ordered in increasing sense and Inf<Index>()-terminated), and the node potential of *nms*[*i*] is written in *P*[*i*]. Note that, unlike [MCFGetX\(\)](#) above, *nms* is an *input* of the method.

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the nodes 'i' with *strt* <= *i* < min([MCFn\(\)](#) , *stp*). 'strt' and 'stp' work in "&&" with *nms*; that is, if *nms* != NULL then only the values corresponding to nodes which are *both* in *nms*[] and whose index is in the correct range are returned.

7.3.5.45 **virtual CNumber MCFClass::MCFGetRC (cIndex *i*) [pure virtual]**

Return the reduced cost of the *i*-th arc. This information should be cheaply available in most implementations.

Note

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

Implemented in [RelaxIV](#).

7.3.5.46 **virtual cCRow MCFClass::MCFGetRC (void) [inline, virtual]**

Return a read-only pointer to an internal data structure containing the reduced costs. Since this may *not* always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Note

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

Reimplemented in [RelaxIV](#).

7.3.5.47 **virtual void MCFClass::MCFGetRC (CRow *CR*, cIndex_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]**

Write the reduced costs corresponding to the current dual solution in *RC*[]. If *nms* == NULL, the reduced cost of arc 'i' (i in 0 .. *m* - 1) is written in *RC*[*i*]; if *nms* != NULL, it must point to a vector of indices in 0 .. *m* - 1 (ordered in increasing sense and Inf<Index>()-terminated), and the reduced cost of arc *nms*[*i*] is written in *RC*[*i*]. Note that, unlike [MCFGetX\(\)](#) above, *nms* is an *input* of the method.

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with *strt* <= *i* < min([MCFm\(\)](#) , *stp*). 'strt' and 'stp' work in "&&" with *nms*; that is, if *nms* != NULL then only the values corresponding to arcs which are *both* in *nms*[] and whose index is in the correct range are returned.

Note

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

7.3.5.48 virtual MCFStatePtr MCFClass::MCFGetState (void) [inline, virtual]

Save the state of the MCF solver. The [MCFClass](#) interface supports the notion of saving and restoring the state of the MCF solver, such as the current/optimal basis in a simplex solver. The "empty" class [MCFState](#) is defined as a placeholder for state descriptions.

[MCFGetState\(\)](#) creates and returns a pointer to an object of (a proper derived class of) class [MCFState](#) which describes the current state of the MCF solver.

Reimplemented in [RelaxIV](#).

7.3.5.49 int MCFClass::MCFGetStatus (void) [inline]

Returns an int describing the current status of the MCF solver. Possible return values are:

- kUnsolved [SolveMCF\(\)](#) has not been called yet, or the data of the problem has been changed since the last call;
- kOK optimization has been carried out successfully;
- kStopped optimization have been stopped before that the stopping conditions of the solver applied, e.g. because of the maximum allowed number of "iterations" [see [SetPar\(int \)](#)] or the maximum allowed time [see [SetPar\(double \)](#)] has been reached; this is not necessarily an error, as it might just be required to re-call [SolveMCF\(\)](#) giving it more "resources" in order to solve the problem;
- kUnfeasible if the current MCF instance is (primal) unfeasible;
- kUnbounded if the current MCF instance is (primal) unbounded (this can only happen if the solver actually allows F_INF capacities, which is nonstandard in the interface);
- kError if there was an error during the optimization; this typically indicates that computation cannot be resumed, although solver-dependent ways of dealing with solver-dependent errors may exist.

[MCFClass](#) has a protected `int` member `status` that can be used by derived classes to hold status information and that is returned by the standard implementation of this method. Note that `status` is an `int` and not an `enum`, and that an `int` is returned by this method, in order to allow the derived classes to extend the set of return values if they need to do so.

7.3.5.50 virtual Index MCFClass::MCFGetUnbCycl (Index_Set Pred, Index_Set ArcPred) [inline, virtual]

Return an unboundedness certificate. In an unbounded MCF problem, unboundedness can always be reduced to the existence of a directed cycle with negative cost and all arcs having infinite capacity. When detecting unboundedness, MCF solvers are typically capable to provide one such cycle. This information can be useful, and this method is provided for getting it. It can be called only if [MCFGetStatus\(\)](#) == kUnbounded, and writes in `Pred[]` and `ArcPred[]`, respectively, the node and arc predecessor function of the cycle. That is, if node 'i' belongs to the cycle then 'Pred[i]' contains the name of the predecessor of 'j' of 'i' in the cycle (note that node names depend on `USENAME0`), and 'ArcPred[i]' contains the index of the arc joining the two (note that in general there may be multiple copies of each arc). Entries of the vectors for nodes not belonging to the cycle are in principle undefined, and the name of one node belonging to the cycle is returned by the method. Note that if there are multiple cycles with negative costs this method will

return just one of them (finding the cycle with most negative cost is an NO-hard problem), although solvers should be able to produce cycles with "large negative" cost.

However, not all solvers may be (easily) capable of providing this information; thus, returning `Inf<Index>()` is allowed, as in the base class implementation, to signify that this information is not available.

7.3.5.51 **virtual FNumber MCFClass::MCFGetUnfCut (Index_Set Cut) [inline, virtual]**

Return an unfeasibility certificate. In an unfeasible MCF problem, unfeasibility can always be reduced to the existence of a cut (subset of nodes of the graph) such as either:

- the inverse of the deficit of the cut (the sum of all the deficits of the nodes in the cut) is larger than the forward capacity of the cut (sum of the capacities of forward arcs in the cut); that is, the nodes in the cut globally produce more flow than can be routed to sinks outside the cut;
- the deficit of the cut is larger than the backward capacity of the cut (sum of the capacities of backward arcs in the cut); that is, the nodes in the cut globally require more flow than can be routed to them from sources outside the cut.

When detecting unfeasibility, MCF solvers are typically capable to provide one such cut. This information can be useful - typically, the only way to make the problem feasible is to increase the capacity of at least one of the forward/backward arcs of the cut -, and this method is provided for getting it. It can be called only if `MCFGetStatus() == kUnfeasible`, and should write in `Cut` the set of names of nodes in the unfeasible cut (note that node names depend on `USENAME0`), `Inf<Index>()`-terminated, returning the deficit of the cut (which allows to distinguish which of the two cases above hold). In general, no special properties can be expected from the returned cut, but solvers should be able to provide e.g. "small" cuts.

However, not all solvers may be (easily) capable of providing this information; thus, returning 0 (no cut) is allowed, as in the base class implementation, to signify that this information is not available.

7.3.5.52 **virtual cFRow MCFClass::MCFGetX (void) [inline, virtual]**

Return a read-only pointer to an internal data structure containing the flow solution in "dense" format. Since this may *not always be available*, depending on the implementation, this method can (uniformly) return `NULL`. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented in [RelaxIV](#).

7.3.5.53 **virtual void MCFClass::MCFGetX (FRow F, Index_Set nms = NULL, cIndex strt = 0, Index stp = Inf< Index >()) [pure virtual]**

Write the optimal flow solution in the vector `F[]`. If `nms == NULL`, `F[]` will be in "dense" format, i.e., the flow relative to arc 'i' (i in 0 .. m - 1) is written in `F[i]`. If `nms != NULL`, `F[]` will be in "sparse" format, i.e., the indices of the nonzero elements in the flow solution are written in `nms` (that is then `Inf<Index>()`-terminated) and the flow value of arc `nms[i]` is written in `F[i]`. Note that `nms` is not* guaranteed to be ordered. Also, note that, unlike `MCFGetRC()` and `MCFGetPi()` [see below], `nms` is an *output* of the method.

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with `strt <= i < min(MCFm(), stp)`.

7.3.5.54 Index MCFClass::MCFm (void) [inline]

Return the number of arcs in the current graph. The implementation of the method in the base class returns the protected fields `m`, which is provided for derived classes to hold this information.

7.3.5.55 Index MCFClass::MCFmmax (void) [inline]

Return the maximum number of arcs for this instance of [MCFClass](#). The implementation of the method in the base class returns the protected fields `mmax`, which is provided for derived classes to hold this information.

7.3.5.56 Index MCFClass::MCFn (void) [inline]

Return the number of nodes in the current graph. The implementation of the method in the base class returns the protected fields `n`, which is provided for derived classes to hold this information.

7.3.5.57 Index MCFClass::MCFnmax (void) [inline]

Return the maximum number of nodes for this instance of [MCFClass](#). The implementation of the method in the base class returns the protected fields `nmax`, which is provided for derived classes to hold this information.

7.3.5.58 virtual void MCFClass::MCFPutState (MCFStatePtr S) [inline, virtual]

Restore the solver to the state in which it was when the state 'S' was created with [MCFGetState\(\)](#) [see above].

The typical use of this method is the following: a MCF problem is solved and the "optimal state" is set aside. Then the problem changes and it is re-solved. Then, the problem has to be changed again to a form that is close to the original one but rather different from the second one (think of a long backtracking in a Branch & Bound) to which the current "state" refers. Then, the old optimal state can be expected to provide a better starting point for reoptimization [see [ReOptimize\(\)](#) below].

Note, however, that the state is only relative to the optimization process, i.e., this operation is meaningless if the data of the problem has changed in the meantime. So, if a state has to be used for speeding up reoptimization, the following has to be done:

- first, the data of the solver is brought back to **exactly** the same as it was at the moment where the state 'S' was created (prior than this operation a `ReOptimize(false)` call is probably advisable);
- then, [MCFPutState\(\)](#) is called (and `ReOptimize(true)` is called);
- only afterwards the data of the problem is changed to the final state and the problem is solved.

A "put state" operation does not "deplete" the state, which can therefore be used more than once. Indeed, a state is constructed inside the solver for each call to [MCFGetState\(\)](#), but the solver never deletes statuses; this has to be done on the outside when they are no longer needed (the solver must be "resistent" to deletion of the state at any moment).

Since not all the MCF solvers reoptimize (efficiently enough to make these operations worth), an "empty" implementation that does nothing is provided by the base class.

Reimplemented in [RelaxIV](#).

7.3.5.59 virtual cRow MCFClass::MCFQCoef(void) [inline, virtual]

Return a read-only pointer to an internal vector containing the arc costs. Since this may **not* always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready (such as "pure linear" MCF solvers that only work with all zero quadratic coefficients) does not need to implement the method.

7.3.5.60 virtual CNumber MCFClass::MCFQCoef(cIndex i) [inline, virtual]

Parameters

i Return the quadratic coefficients of the cost of the *i*-th arc. Note that the method is **not** pure virtual: an implementation is provided for "pure linear" MCF solvers that only work with all zero quadratic coefficients.

7.3.5.61 virtual void MCFClass::MCFQCoef(CRow Qv, cIndex_Set nms = NULL, cIndex strt = 0, Index stp = Inf<Index>()) [inline, virtual]

Parameters

stp Write the quadratic coefficients of the arc costs into Qv[]. If nms == NULL, then all the costs are written, otherwise Costv[i] contains the information relative to arc nms[i] (nms[] must be Inf<Index>()-terminated).

The parameters 'strt' and 'stp' allow to restrict the output of the method to all and only the arcs 'i' with $\text{strt} \leq i < \min(\text{MCFm}(), \text{stp})$. 'strt' and 'stp' work in "&&" with nms; that is, if nms != NULL then only the values corresponding to arcs which are **both** in nms and whose index is in the correct range are returned.

Note that the method is **not** pure virtual: an implementation is provided for "pure linear" MCF solvers that only work with all zero quadratic coefficients.

7.3.5.62 virtual Index MCFClass::MCFSNde(cIndex i) [pure virtual]

Return the starting (tail) node of the arc 'i'.

Note

If USERNAME0 == 0 then the returned node names will be in the range 1 .. n, while if USERNAME0 == 1 the returned node names will be in the range 0 .. n - 1.

Implemented in [RelaxIV](#).

7.3.5.63 virtual cIndex_Set MCFClass::MCFSNdes(void) [inline, virtual]

Return a read-only pointer to an internal vector containing the starting (tail) nodes for each arc. Since this may **not* always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented in [RelaxIV](#).

7.3.5.64 virtual FNumber MCFClass::MCFUCap (cIndex *i*) [pure virtual]

Return the capacity of the *i*-th arc.

Implemented in [RelaxIV](#).

7.3.5.65 virtual cFRow MCFClass::MCFUCaps (void) [inline, virtual]

Return a read-only pointer to an internal vector containing the arc capacities. Since this may *not* always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented in [RelaxIV](#).

7.3.5.66 virtual void MCFClass::MCFUCaps (FRow *UCapv*, cIndex_Set *nms* = NULL, cIndex *strt* = 0, Index *stp* = Inf< Index >()) [pure virtual]

Write the arc capacities into *UCapv*[]. If *nms* == NULL, then all the capacities are written, otherwise *UCapv*[*i*] contains the information relative to arc *nms*[*i*] (*nms*[] must be Inf<Index>()-terminated).

The parameters '*strt*' and '*stp*' allow to restrict the output of the method to all and only the arcs '*i*' with *strt* ≤ *i* < min([MCFm\(\)](#) , *stp*). '*strt*' and '*stp*' work in "&&" with *nms*; that is, if *nms* != NULL then only the values corresponding to arcs which are *both* in *nms* and whose index is in the correct range are returned.

7.3.5.67 virtual void MCFClass::OpenArc (cIndex *name*) [pure virtual]

Restore the previously closed arc '*name*'. It is an error to open an arc that has not been previously closed.

Implemented in [RelaxIV](#).

7.3.5.68 virtual void MCFClass::PreProcess (void) [inline, virtual]

Extract a smaller/easier equivalent MCF problem. The data of the instance is changed and the easier one is solved instead of the original one. In the MCF case, preprocessing may involve reducing bounds, identifying disconnected components of the graph etc. However, preprocessing is solver-specific.

This method can be implemented by derived classes in their solver-specific way. Preprocessing may reveal unboundedness or unfeasibility of the problem; if that happens, [PreProcess\(\)](#) should properly set the '*status*' field, that can then be read with [MCFGetStatus\(\)](#) [see below].

Note that preprocessing may destroy all the solution information. Also, it may be allowed to change the data of the problem, such as costs/capacities of the arcs.

A valid preprocessing is doing nothing, and that's what the default implementation of this method (that is *not* pure virtual) does.

Reimplemented in [RelaxIV](#).

7.3.5.69 virtual void MCFClass::SetMCFTIME (bool *TimeIt* = true) [inline, virtual]

Allocate an OPTtimers object [see OPTtypes.h] to be used for timing the methods of the class. The time can be read with [TimeMCF\(\)](#) [see below]. By default, or if SetMCFTIME(false) is called, no timing is

done. Note that, since all the relevant methods of the class are pure virtual, `MCFClass` can only manage the `OPTtimers` object, but it is due to derived classes to actually implement the timing.

Note

time accumulates over the calls: calling `SetMCFTime()`, however, resets the counters, allowing to time specific groups of calls.
 of course, setting `kMaxTime` [see `SetPar()` above] to any nonzero value has no effect unless `SetMCFTime(true)` has been called.

7.3.5.70 void MCFClass::SetPar (int par, double val) [inline, virtual]

Set float parameters of the algorithm.

Parameters

par is the parameter to be set; the enum `MCFParam` can be used, but 'par' is an int (every enum is an int) so that the method can be extended by derived classes for the setting of their parameters
value is the value to assign to the parameter.

The base class implementation handles these parameters:

- `kEpsFlw`: sets the tolerance for controlling if the flow on an arc is zero to val. This also sets the tolerance for controlling if a node deficit is zero (see `kEpsDfct`) to $val * < \text{max number of nodes} >$; this value should be safe for graphs in which any node has less than $< \text{max number of nodes} >$ adjacent nodes, i.e., for all graphs but for very dense ones with "parallel arcs"
- `kEpsDfct`: sets the tolerance for controlling if a node deficit is zero to val, in case a better value than that automatically set by `kEpsFlw` (see above) is available (e.g., $val * k$ would be good if no node has more than k neighbours)
- `kEpsCst`: sets the tolerance for controlling if the reduced cost of an arc is zero to val. A feasible solution satisfying eps-complementary slackness, i.e., such that

$$RC[i, j] < -eps \Rightarrow X[i, j] = U[ij]$$

and

$$RC[i, j] > eps \Rightarrow X[i, j] == 0,$$

is known to be $(eps * n)$ -optimal.

- `kMaxTime`: sets the max time (in seconds) in which the MCF Solver can find an optimal solution (default 0, which means no limit).

7.3.5.71 void MCFClass::SetPar (int par, int val) [inline, virtual]

Set integer parameters of the algorithm.

Parameters

par is the parameter to be set; the enum `MCFParam` can be used, but 'par' is an int (every enum is an int) so that the method can be extended by derived classes for the setting of their parameters

value is the value to assign to the parameter.

The base class implementation handles these parameters:

- `kMaxIter`: the max number of iterations in which the MCF Solver can find an optimal solution (default 0, which means no limit)
- `kReopt`: tells the solver if it has to reoptimize. The implementation in the base class sets a flag, the protected `bool` field `Senstv`; if true (default) this field instructs the MCF solver to try to exploit the information about the latest optimal solution to speedup the optimization of the current problem, while if the field is false the MCF solver should restart the optimization "from scratch" discarding any previous information. Usually reoptimization speeds up the computation considerably, but this is not always true, especially if the data of the problem changes a lot.

Reimplemented in [RelaxIV](#).

7.3.5.72 `virtual void MCFClass::SolveMCF(void) [pure virtual]`

Solver of the Min Cost Flow Problem. Attempts to solve the MCF instance currently loaded in the object.

Implemented in [RelaxIV](#).

7.3.5.73 `double MCFClass::TimeMCF(void) [inline]`

Like `TimeMCF(double,double)` [see above], but returns the total time.

7.3.5.74 `void MCFClass::TimeMCF(double & t_us, double & t_ss) [inline]`

Time the code. If called within any of the methods of the class that are "actively timed" (this depends on the subclasses), this method returns the user and sistem time (in seconds) since the start of that method. If methods that are actively timed call other methods that are actively timed, [TimeMCF\(\)](#) returns the (...) time since the beginning of the outer* actively timed method. If called outside of any actively timed method, this method returns the (...) time spent in all the previous executions of all the actively timed methods of the class.

Implementing the proper calls to `MCFt->Start()` and `MCFt->Stop()` is due to derived classes; these should at least be placed at the beginning and at the end, respectively, of [SolveMCF\(\)](#) and presumably the `Chg***()` methods, that is, at least these methods should be "actively timed".

7.3.5.75 `void MCFClass::WriteMCF(ostream & oStrm, int frmt = 0) [inline, virtual]`

Write the current MCF problem to an ostream. This may be useful e.g. for debugging purposes.

The base [MCFClass](#) class provides output in two different formats, depending on the value of the parameter `frmt`:

- `kDimacs` the problem is written in DIMACS standard format, read by most MCF codes available;
- `kMPS` the problem is written in the "modern version" (tab-separated) of the MPS format, read by most LP/MIP solvers;

- kFWMPS the problem is written in the "old version" (fixed width fields) of the MPS format; this is read by most LP/MIP solvers, but some codes still require the old format.

The implementation of [WriteMCF\(\)](#) in the base class uses all the above methods for reading the data; as such it will work for any derived class that properly implements this part of the interface, but it may not be very efficient. Thus, the method is virtual to allow the derived classes to either re-implement [WriteMCF\(\)](#) for the above two formats in a more efficient way, and/or to extend it to support other solver-specific formats.

Note

None of the above two formats supports quadratic MCFs, so if nonzero quadratic coefficients are present, they are just ignored.

Reimplemented in [RelaxIV](#).

7.3.6 Member Data Documentation

7.3.6.1 `int MCFClass::MaxIter` `[protected]`

max number of iterations in which MCF Solver can find an optimal solution (0 = no limits)

7.3.6.2 `double MCFClass::MaxTime` `[protected]`

max time (in seconds) in which MCF Solver can find an optimal solution (0 = no limits)

7.3.6.3 `int MCFClass::status` `[protected]`

return status, see the comments to [MCFGetStatus\(\)](#) above. Note that the variable is defined int to allow derived classes to return their own specialized status codes

The documentation for this class was generated from the following file:

- [MCFClass.h](#)

7.4 MCFClass::MCFException Class Reference

```
#include <MCFClass.h>
```

Public Member Functions

- **MCFException** (const char *const msg=0)
- const char * **what** (void) const throw ()

7.4.1 Detailed Description

Small class for exceptions. Derives from std::exception implementing the virtual method what() -- and since what is virtual, remember to always catch it by reference (catch exception &e) if you want the thing to work. [MCFException](#) class are thought to be of the "fatal" type, i.e., problems for which no solutions exists apart from aborting the program. Other kinds of exceptions can be properly handled by defining derived classes with more information.

The documentation for this class was generated from the following file:

- [MCFClass.h](#)

7.5 MCFClass::MCFState Class Reference

```
#include <MCFClass.h>
```

Inherited by RelaxIV::RIVState.

7.5.1 Detailed Description

Base class for representing the internal state of the MCF algorithm.

The documentation for this class was generated from the following file:

- [MCFClass.h](#)

7.6 OPTrand Class Reference

```
#include <OPTUtils.h>
```

Public Member Functions

- [OPTrand](#) (void)
constructor of the class
- double [rand](#) (void)
- void [srand](#) (long seed)
Seeds the random generator for this instance of [OPTrand](#).

7.6.1 Detailed Description

Provide a common interface to the different random generators that are available in different platforms.

7.6.2 Member Function Documentation

7.6.2.1 double OPTrand::rand (void) [inline]

Returns a random number uniformly distributed in [0, 1).

Note

each object of class [OPTrand](#) has its own sequence, so that multiple [OPTrand](#) objects being used within the same program do not interfere with each other (as opposed to what C random routines would do).


The documentation for this class was generated from the following file:

- [OPTUtils.h](#)

7.7 RelaxIV Class Reference

```
#include <RelaxIV.h>
```

Inheritance diagram for RelaxIV:



Classes

- class **RIVState**

Public Types

- enum **MCFRParam** { **kAuction** = kLastParam }
- enum **RIVFIFmt** { **kCLP** = kMPS + 1, **kRIV** }
- typedef bool * **Bool_Vec**
vector of booleans

Public Member Functions

- **RelaxIV** (**cIndex** nmX=0, **cIndex** mmX=0)
- void **LoadNet** (**cIndex** nmX=0, **cIndex** mmX=0, **cIndex** pn=0, **cIndex** pm=0, **cFRow** pU=NULL, **cCRow** pC=NULL, **cFRow** pDfct=NULL, **cIndex_Set** pSn=NULL, **cIndex_Set** pEn=NULL)
- virtual void **SetPar** (int par, int val)
- virtual void **GetPar** (int par, int &val)
- void **PreProcess** (void)
- void **SolveMCF** (void)
- void **MCFGetX** (register **FRow** F, register **Index_Set** nms=NULL, **cIndex** strt=0, **Index** stp=Inf<**Index**>())
- **cFRow** **MCFGetX** (void)
- void **MCFGetRC** (register **CRow** CR, register **cIndex_Set** nms=NULL, **cIndex** strt=0, **Index** stp=Inf<**Index**>())
- **cCRow** **MCFGetRC** (void)
- **CNumber** **MCFGetRC** (**cIndex** i)
- void **MCFGetPi** (register **CRow** P, register **cIndex_Set** nms=NULL, **cIndex** strt=0, **Index** stp=Inf<**Index**>())
- **cCRow** **MCFGetPi** (void)
- **RelaxIV::FONumber** **MCFGetFO** (void)
- **MCFClass::MCFStatePtr** **MCFGetState** (void)
- void **MCFPutState** (**MCFClass::MCFStatePtr** S)
- void **MCFArcs** (register **Index_Set** Startv, register **Index_Set** Endv, register **cIndex_Set** nms=NULL, **cIndex** strt=0, **Index** stp=Inf<**Index**>())
- **Index** **MCFNde** (**cIndex** i)

- [Index](#) [MCFENde](#) ([cIndex](#) i)
- [cIndex_Set](#) [MCFSNdes](#) (void)
- [cIndex_Set](#) [MCFENdes](#) (void)
- void [MCFCosts](#) (register [CRow](#) Costv, register [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=[Inf](#)< [Index](#) >())
- [CNumber](#) [MCFCost](#) ([cIndex](#) i)
- [cCRow](#) [MCFCosts](#) (void)
- void [MCFUCaps](#) (register [FRow](#) UCapv, register [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=[Inf](#)< [Index](#) >())
- [FNumber](#) [MCFUCap](#) ([cIndex](#) i)
- [cFRow](#) [MCFUCaps](#) (void)
- void [MCFDfcts](#) (register [FRow](#) Dfctv, register [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=[Inf](#)< [Index](#) >())
- [FNumber](#) [MCFDfct](#) ([cIndex](#) i)
- [cFRow](#) [MCFDfcts](#) (void)
- void [WriteMCF](#) (ostream &oStrm, int frmt=0)
- void [ChgCosts](#) (register [cCRow](#) NCost, register [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=[Inf](#)< [Index](#) >())
- void [ChgCost](#) (register [Index](#) arc, [cCNumber](#) NCost)
- void [ChgDfcts](#) (register [cFRow](#) NDfct, register [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=[Inf](#)< [Index](#) >())
- void [ChgDfct](#) (register [Index](#) nod, [cFNumber](#) NDfct)
- void [ChgUCaps](#) (register [cFRow](#) NCap, register [cIndex_Set](#) nms=NULL, [cIndex](#) strt=0, [Index](#) stp=[Inf](#)< [Index](#) >())
- void [ChgUCap](#) (register [Index](#) arc, [cFNumber](#) NCap)
- void [CloseArc](#) ([cIndex](#) name)
- bool [IsClosedArc](#) ([cIndex](#) name)
- void [DelNode](#) ([cIndex](#) name)
- void [OpenArc](#) ([cIndex](#) name)
- [Index](#) [AddNode](#) ([cFNumber](#) aDfct)
- void [ChangeArc](#) ([cIndex](#) name, [cIndex](#) nSS=[Inf](#)< [Index](#) >(), [cIndex](#) nEN=[Inf](#)< [Index](#) >())
- void [DelArc](#) ([cIndex](#) name)
- bool [IsDeletedArc](#) ([cIndex](#) name)
- [Index](#) [AddArc](#) ([cIndex](#) Start, [cIndex](#) End, [cFNumber](#) aU, [cCNumber](#) aC)
- int [MCFiter](#) ()

Total number of (single-node or multinode) iterations.

- int [MCFaug](#) ()

Number of flow augmentations.

7.7.1 Detailed Description

The [RelaxIV](#) class derives from the abstract base class [MCFClass](#), thus sharing its (standard) interface, and implements a Relaxation algorithm for solving (Linear) Min Cost Flow problems.

7.7.2 Member Enumeration Documentation

7.7.2.1 enum RelaxIV::MCFRParam

Public enum describing the possible parameters of the MCF solver, "extended" from [MCFCClass::MCFParam](#), to be used with the methods [SetPar\(\)](#) and [GetPar\(\)](#).

Enumerator:

kAuction crash initialization

7.7.2.2 enum RelaxIV::RIVFIrmt

Public enum describing the more file formats in [RelaxIV::WriteMCF\(\)](#).

Enumerator:

kCLP the "LP" format

kRIV RelaxIV-specific format.

7.7.3 Constructor & Destructor Documentation

7.7.3.1 RelaxIV::RelaxIV (cIndex *nm*x = 0, cIndex *mm*x = 0)

Constructor of the class, as in [MCFCClass::MCFCClass\(\)](#).

7.7.4 Member Function Documentation

7.7.4.1 MCFCClass::Index RelaxIV::AddArc (cIndex *Start*, cIndex *End*, cFNumber *aU*, cCNumber *aC*) [**virtual**]

Add the new arc (*Start* , *End*) with cost *aC* and capacity *aU*, returning its name. [Inf<Index>\(\)](#) is returned if there is no room for a new arc. Remember that arc names go from 0 to *mm*x - 1.

Implements [MCFCClass](#).

7.7.4.2 MCFCClass::Index RelaxIV::AddNode (cFNumber *aDfct*) [**virtual**]

Add a new node with deficit *aDfct*, returning its name. [Inf<Index>\(\)](#) is returned if there is no room for a new node. Remember that the node names are either { 0 .. *n*max - 1 } or { 1 .. *n*max }, depending on the value of `USENAME0`.

Implements [MCFCClass](#).

7.7.4.3 void RelaxIV::ChangeArc (cIndex *name*, cIndex *nSN* = [Inf< Index >\(\)](#), cIndex *nEN* = [Inf< Index >\(\)](#)) [**virtual**]

Change the starting and/or ending node of arc 'name' to *nSN* and *nEN*. Each parameter being [Inf<Index>\(\)](#) means to leave the previous starting or ending node untouched. When this method is called 'name' can be either the name of a "normal" arc or that of a "closed" arc [see [CloseArc\(\)](#) above]: in the latter case, at the end of [ChangeArc\(\)](#) the arc is *still closed*, and it remains so until [OpenArc\(name \)](#) [see above] is called.

Implements [MCFCClass](#).

7.7.4.4 void RelaxIV::CloseArc (cIndex *name*) [virtual]

"Close" the arc '*name*'. Although all the associated information (name, cost, capacity, end and start node) is kept, the arc is removed from the problem until `OpenArc(i)` [see below] is called.

"closed" arcs always have 0 flow, but are otherwise counted as any other arc; for instance, `MCFm()` does *not* decrease as an effect of a call to `CloseArc()`. How this closure is implemented is solver-specific.

Implements [MCFClass](#).

7.7.4.5 void RelaxIV::DelArc (cIndex *name*) [virtual]

Delete the arc '*name*'. Unlike "closed" arcs, all the information associated with a deleted arc is lost and '*name*' is made available as a name for new arcs to be created with `AddArc()` [see below].

If furthermore '*name*' is the last arc, the number of arcs as reported by `MCFm()` is reduced by at least one, until the *m*-th arc is not a deleted one. Otherwise, the flow on the arc is always ensured to be 0.

Implements [MCFClass](#).

7.7.4.6 void RelaxIV::DelNode (cIndex *name*) [virtual]

Delete the node '*name*'.

For any value of '*name*', all incident arcs to that node are closed [see `CloseArc()` above] (*not* Deleted, see `DelArc()` below) and the deficit is set to zero.

If furthermore '*name*' is the last node, the number of nodes as reported by `MCFn()` is reduced by at least one, until the *n*-th node is not a deleted one.

Implements [MCFClass](#).

7.7.4.7 void RelaxIV::GetPar (int *par*, int & *val*) [inline, virtual]

This method returns one of the integer parameter of the algorithm.

Parameters

par is the parameter to return [see `SetPar(int)` for comments];

val upon return, it will contain the value of the parameter.

Apart from the parameters of the base class, this method handles `kAuction`.

Reimplemented from [MCFClass](#).

7.7.4.8 bool RelaxIV::IsClosedArc (cIndex *name*) [inline, virtual]

`IsClosedArc()` returns true if and only if the arc '*name*' is closed.

Implements [MCFClass](#).

7.7.4.9 bool RelaxIV::IsDeletedArc (cIndex *name*) [inline, virtual]

Return true if and only if the arc '*name*' is deleted. It should only be called with `name < MCFm()`, as every other arc is deleted by definition.

Implements [MCFClass](#).

7.7.4.10 void RelaxIV::LoadNet (cIndex *nm*x = 0, cIndex *mm*x = 0, cIndex *pn* = 0, cIndex *pm* = 0, cFRow *pU* = NULL, cCRow *pC* = NULL, cFRow *pDfct* = NULL, cIndex_Set *pSn* = NULL, cIndex_Set *pEn* = NULL) [virtual]

Inputs a new network, as in [MCFClass::LoadNet\(\)](#).

Arcs with $pC[i] == \text{Inf} < \text{CNumber} >()$ do not "exist". If $\text{DYNMC_MCF_RIV} > 0$, these arcs are "closed".

If $\text{DYNMC_MCF_RIV} == 0$ but $\text{SAME_GRPH_RIV} > 0$, these arcs are dealt with exactly like as if they had $pU[i] == 0$, i.e., as "normal" arcs with zero capacity. These arcs can be put back into the formulation by simply changing their capacity (and cost). Note that, however, this is less efficient than eliminating them explicitly from the problem.

If $\text{DYNMC_MCF_RIV} == 0$ and $\text{SAME_GRPH_RIV} == 0$, these arcs are just removed from the formulation. However, they have some sort of a "special status" (after all, if the user wants to remove them completely he/she can just change the data), in that they are still counted into the number of arcs of the graph and they will always have 0 flow and $\text{Inf} < \text{CNumber} >()$ reduced cost as "closed" or "deleted" arcs.

If $\text{SAME_GRPH_RIV} == 1$ ($\Rightarrow \text{DYNMC_MCF_RIV} < 3$), *pSn* and *pEn* passed to any instance after the first that is constructed are ignored, thus, $pSn == pEn == \text{NULL}$ is allowed in this case. Note that *pn* and *pm* are not* ignored.

Implements [MCFClass](#).

7.7.4.11 MCFClass::CNumber RelaxIV::MCFCost (cIndex *i*) [inline, virtual]

Return the cost of the *i*-th arc.

Implements [MCFClass](#).

7.7.4.12 MCFClass::cCRow RelaxIV::MCFCosts (void) [inline, virtual]

Return a read-only pointer to an internal vector containing the arc costs. Since this may *not* always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented from [MCFClass](#).

7.7.4.13 MCFClass::FNumber RelaxIV::MCFDfct (cIndex *i*) [inline, virtual]

Return the deficit of the *i*-th node.

Note

Here node "names" go from 0 to *n* - 1, regardless to the value of `USENAME0`; hence, if `USENAME0 == 0` then the first node is "named 1", but its deficit is returned by `MCFDfct(0)`.

Implements [MCFClass](#).

7.7.4.14 MCFClass::cFRow RelaxIV::MCFDfcts (void) [inline, virtual]

Return a read-only pointer to an internal vector containing the node deficits. Since this may *not always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Note

Here node "names" go from 0 to n - 1, regardless to the value of USERNAME0; hence, if USERNAME0 == 0 then the first node is "named 1", but its deficit is contained in [MCFDfcts\(\)\[0 \]](#).

Reimplemented from [MCFClass](#).

7.7.4.15 MCFClass::Index RelaxIV::MCFENde (cIndex i) [inline, virtual]

Return the ending (head) node of the arc 'i'.

Note

If USERNAME0 == 0 then the returned node names will be in the range 1 .. n, while if USERNAME0 == 1 the returned node names will be in the range 0 .. n - 1.

Implements [MCFClass](#).

7.7.4.16 MCFClass::cIndex_Set RelaxIV::MCFENdes (void) [inline, virtual]

Same meaning as [MCFClass::MCFENdes\(\)](#).

Note

[MCFENdes\(\)](#) returns a pointers to a (read-only) vector containing the arc end nodes *only if USERNAME0 == 0*; otherwise, it returns NULL.

Reimplemented from [MCFClass](#).

7.7.4.17 MCFClass::FONumber RelaxIV::MCFGetFO (void) [virtual]

Return the objective function value of the primal solution currently returned by [MCFGetX\(\)](#).

If [MCFGetStatus\(\)](#) == kOK, this is guaranteed to be the optimal objective function value of the problem (to within the optimality tolerances), but only prior to any call to [HaveNewX\(\)](#) that returns true. [MCFGetFO\(\)](#) typically returns Inf<FONumber>() if [MCFGetStatus\(\)](#) == kUnfeasible and

- Inf<FONumber>() if [MCFGetStatus\(\)](#) == kUnbounded. If [MCFGetStatus\(\)](#) == kStopped and [MCFGetFO\(\)](#) returns a finite value, it must be an upper bound on the optimal objective function value (typically, the objective function value of one primal feasible solution).

Implements [MCFClass](#).

7.7.4.18 MCFClass::cCRow RelaxIV::MCFGetPi (void) [virtual]

Same meaning as [MCFClass::MCFGetPi\(\)](#).

Note

if [MCFGetPi\(void \)](#) returns a pointer, this is a pointer to a static shared data structure that can be corrupted by calls to methods (e.g. [MCFGetPi\(\)](#) itself) for other active instances of [RelaxIV](#).

Reimplemented from [MCFClass](#).

7.7.4.19 void RelaxIV::MCFGetPi (register CRow P, register cIndex_Set nms = NULL, cIndex strt = 0, Index stp = Inf< Index > ())

Same meaning as [MCFClass::MCFGetPi\(\)](#).

Note

if both AUCTION == 0 and DYNMC_MCF_RIV <= 1, no internal memory for the vector of potentials is allocated; hence, even if nms != NULL MCFGetPi(P) first constructs the full vector of potentials in P and then selects only the components in nms. Therefore, memory can be written even after the | nms | -th element of P. For the same reason, memory can be written even after the (stp - strt) -th element of P.

7.7.4.20 MCFClass::CNumber RelaxIV::MCFGetRC (cIndex i) [virtual]

Return the reduced cost of the i-th arc. This information should be cheaply available in most implementations.

Note

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

Implements [MCFClass](#).

7.7.4.21 MCFClass::cCRow RelaxIV::MCFGetRC (void) [virtual]

Return a read-only pointer to an internal data structure containing the reduced costs. Since this may *not* always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Note

the output of [MCFGetRC\(\)](#) will change after any call to [HaveNewPi\(\)](#) [see above] which returns true.

Reimplemented from [MCFClass](#).

7.7.4.22 MCFClass::MCFStatePtr RelaxIV::MCFGetState (void) [virtual]

Same meaning as [MCFClass::MCFGetState\(\)](#).

The state of the algorithm is the pair $S = (X[] , RC[])$ of the arc flows and reduced costs.

Reimplemented from [MCFClass](#).

7.7.4.23 MCFClass::cFRow RelaxIV::MCFGetX (void) [virtual]

Return a read-only pointer to an internal data structure containing the flow solution in "dense" format. Since this may *not always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented from [MCFClass](#).

7.7.4.24 void RelaxIV::MCFPutState (MCFClass::MCFStatePtr S) [virtual]

Restore the solver to the state in which it was when the state 'S' was created with [MCFGetState\(\)](#) [see above].

The typical use of this method is the following: a MCF problem is solved and the "optimal state" is set aside. Then the problem changes and it is re-solved. Then, the problem has to be changed again to a form that is close to the original one but rather different from the second one (think of a long backtracking in a Branch & Bound) to which the current "state" refers. Then, the old optimal state can be expected to provide a better starting point for reoptimization [see [ReOptimize\(\)](#) below].

Note, however, that the state is only relative to the optimization process, i.e., this operation is meaningless if the data of the problem has changed in the meantime. So, if a state has to be used for speeding up reoptimization, the following has to be done:

- first, the data of the solver is brought back to *exactly* the same as it was at the moment where the state 'S' was created (prior than this operation a [ReOptimize\(false \)](#) call is probably advisable);
- then, [MCFPutState\(\)](#) is called (and [ReOptimize\(true \)](#) is called);
- only afterwards the data of the problem is changed to the final state and the problem is solved.

A "put state" operation does not "deplete" the state, which can therefore be used more than once. Indeed, a state is constructed inside the solver for each call to [MCFGetState\(\)](#), but the solver never deletes statuses; this has to be done on the outside when they are no longer needed (the solver must be "resistent" to deletion of the state at any moment).

Since not all the MCF solvers reoptimize (efficiently enough to make these operations worth), an "empty" implementation that does nothing is provided by the base class.

Reimplemented from [MCFClass](#).

7.7.4.25 MCFClass::Index RelaxIV::MCFSNde (cIndex i) [inline, virtual]

Return the starting (tail) node of the arc 'i'.

Note

If `USENAME0 == 0` then the returned node names will be in the range 1 .. n, while if `USENAME0 == 1` the returned node names will be in the range 0 .. n - 1.

Implements [MCFClass](#).

7.7.4.26 MCFClass::cIndex_Set RelaxIV::MCFSNdes (void) [inline, virtual]

Same meaning as [MCFClass::MCFSNdes\(\)](#).

Note

[MCFSNdes\(\)](#) returns a pointers to a (read-only) vector containing the arc start nodes *only if USE-NAME0 == 0*; otherwise, it returns NULL.

Reimplemented from [MCFClass](#).

7.7.4.27 MCFClass::FNumber RelaxIV::MCFUCap (cIndex i) [inline, virtual]

Return the capacity of the i-th arc.

Implements [MCFClass](#).

7.7.4.28 MCFClass::cFRow RelaxIV::MCFUCaps (void) [inline, virtual]

Return a read-only pointer to an internal vector containing the arc capacities. Since this may *not always be available*, depending on the implementation, this method can (uniformly) return NULL. This is done by the base class already, so a derived class that does not have the information ready does not need to implement the method.

Reimplemented from [MCFClass](#).

7.7.4.29 void RelaxIV::OpenArc (cIndex name) [virtual]

Restore the previously closed arc 'name'. It is an error to open an arc that has not been previously closed.

Implements [MCFClass](#).

7.7.4.30 void RelaxIV::PreProcess (void) [virtual]

If this method is called, a preprocessing phase is performed trying to reduce the arc capacities. This may sometimes help in speeding up the solution of the problem, but may also change the capacities returned by [MCFUCap\[s\]\(\)](#) [see below].

For this method to work properly, arc capacities, node deficits and the topology of the graph must have already been provided with [LoadNet\(\)](#) [see above].

This method can be called more than once, for instance whenever the capacities of some arcs or the deficits of some nodes are changed; however, it destroys the previous optimal solution (if any), forcing the algorithm to restart from scratch.

Reimplemented from [MCFClass](#).

7.7.4.31 void RelaxIV::SetPar (int par, int val) [inline, virtual]

Set integer parameters of the algorithm.

Parameters

par is the parameter to be set;

value is the value to assign to the parameter.

Apart from the parameters of the base class, this method handles:

- `kAuction`: if set to `kYes`, the auction/shortest paths initialization is used in [SolveMCF\(\)](#) to generate the starting solution; if set to `kNo` (default), then the default initialization based on special single-node relaxation iterations is used instead. Note that this parameter is **ignored** if `AUCTION == 0`.

Reimplemented from [MCFClass](#).

7.7.4.32 void RelaxIV::SolveMCF (void) [virtual]

Solver of the Min Cost Flow Problem. Attempts to solve the MCF instance currently loaded in the object.

Implements [MCFClass](#).

7.7.4.33 void RelaxIV::WriteMCF (ostream & oStrm, int frmt = 0) [virtual]

Write the current MCF problem to an ostream. This may be useful e.g. for debugging purposes.

The base [MCFClass](#) class provides output in two different formats, depending on the value of the parameter `frmt`:

- `kDimacs` the problem is written in DIMACS standard format, read by most MCF codes available;
- `kMPS` the problem is written in the "modern version" (tab-separated) of the MPS format, read by most LP/MIP solvers;
- `kFWMPS` the problem is written in the "old version" (fixed width fields) of the MPS format; this is read by most LP/MIP solvers, but some codes still require the old format.

The implementation of [WriteMCF\(\)](#) in the base class uses all the above methods for reading the data; as such it will work for any derived class that properly implements this part of the interface, but it may not be very efficient. Thus, the method is virtual to allow the derived classes to either re-implement [WriteMCF\(\)](#) for the above two formats in a more efficient way, and/or to extend it to support other solver-specific formats.

Note

None of the above two formats supports quadratic MCFs, so if nonzero quadratic coefficients are present, they are just ignored.

Reimplemented from [MCFClass](#).

The documentation for this class was generated from the following files:

- [RelaxIV.h](#)
- [RelaxIV.C](#)

Chapter 8

File Documentation

8.1 Main.C File Reference

```
#include <fstream>
#include <sstream>
#include "RelaxIV.h"
#include <stdlib.h>
#include <time.h>
#include <iostream>
#include <iomanip>
```

Defines

- #define **MCFSOLVER** [RelaxIV](#)
- #define **PRINT_RESULTS** 0

Functions

- template<class T >
 T **ABS** (const T x)
- void **SkipComments** (ifstream &iParam, string &buf)
- void **SetParam** ([MCFCClass](#) *mcf)
- int **main** (int argc, char **argv)

8.1.1 Detailed Description

Sample Main file to illustrate the use of any solver deriving from [MCFCClass](#). By changing just *two lines of code* and little more (see comment PECULIARITY, if exists) the file works with any derived solver.

An instance of a Min Cost Flow problem in DIMACS standard format is read from file and solved. In addition, the same problem can be written on a file in MPS format.

Version

4.00

Date

30 - 12 - 2009

Author

Alessandro Bertolini
Operations Research Group
Dipartimento di Informatica
Universita' di Pisa
Antonio Frangioni
Operations Research Group
Dipartimento di Informatica
Universita' di Pisa

/*-----

8.2 MCFCClass.h File Reference

```
#include "OPTUtils.h"  
#include <iomanip>  
#include <sstream>
```

Classes

- class [MCFCClass](#)
- class [MCFCClass::Inf< T >](#)
- class [MCFCClass::Eps< T >](#)
- class [MCFCClass::MCFCException](#)
- class [MCFCClass::MCFState](#)

Defines

- #define [USENAME0](#) 0

8.2.1 Detailed Description

Header file for the abstract base class [MCFCClass](#), which defines a standard interface for (linear or convex quadratic separable) Min Cost Flow Problem solvers, to be implemented as derived classes.

Version

3.00

Date

30 - 12 - 2009

Author

Alessandro Bertolini
Operations Research Group
Dipartimento di Informatica
Universita' di Pisa
Antonio Frangioni
Operations Research Group
Dipartimento di Informatica
Universita' di Pisa
Claudio Gentile
Istituto di Analisi di Sistemi e Informatica
Consiglio Nazionale delle Ricerche

Copyright © 1996 - 2009 by Antonio Frangioni, Claudio Gentile

8.3 OPTUtils.h File Reference

```
#include <sys/timeb.h>
#include <limits.h>
#include <iostream>
```

Classes

- class [OPTrand](#)

Defines

- #define [OPT_USE_NAMESPACES](#) 0
- #define [OPT_TIMERS](#) 5
- #define [OPT_RANDOM](#) 1
- #define [NULL](#) 0
- #define [CLK_TCK](#) 100

Functions

- template<class T >
void [DfltDStInpt](#) (istream *iStrm, T &Param, const T Dflt, const char cmntc= '#')

8.3.1 Detailed Description

Small classes are provided for:

- reading the time of a code;
- generating random numbers.

The classes can be adapted to different environments setting a compile-time switch in this file.

Additionally, a function is provided for safely reading parameters out of a stream.

Version

2.63

Date

29 - 08 - 2005

Author

Antonio Frangioni
Operations Research Group
Dipartimento di Informatica
Universita' di Pisa

Copyright 1994 - 2004

8.4 RelaxIV.h File Reference

```
#include "MCFClass.h"
```

Classes

- class [RelaxIV](#)
- class [RelaxIV::RIVState](#)

Defines

- #define [SAME_GRPH_RIV](#) 0
- #define [DYNMC_MCF_RIV](#) 0
- #define [AUCTION](#) 0
- #define [RELAXIV_STATISTICS](#) 0

8.4.1 Detailed Description

Linear Min Cost Flow problems solver, based on the [RelaxIV](#) code by D. Bertsekas. Conforms to the standard (MCF) interface defined in [MCFClass.h](#).

The [RelaxIV](#) is based on a primal-dual algorithm which essentially operates as follows: a pseudoflow (a flow vector which satisfies bound and non-negativity constraints but not necessarily flow conservation constraints) is kept which satisfies complementarity slackness conditions with the current vector of potentials; that is, only the flow on arcs whose reduced cost

$$RC[i, j] = C[i, j] - P_i[j] + P_i[i]$$

is zero can be chosen to any value between 0 and the capacity, while arcs with reduced cost < 0 are saturated (fixed to their capacity) and arcs with reduced cost > 0 are empty (fixed to 0).

The algorithm attempts to convert the pseudoflow into a flow (i.e., to satisfy the flow conservation constraints) by essentially running a max-flow algorithm of the augmenting path type. If the flow is found then this is an optimal solution of the problem and the algorithm is stopped. Otherwise, a saturated cut is identified which separates the origins (nodes not yet producing enough flow) to the destinations (nodes not yet consuming enough flow); this cut is used to modify the potentials, thereby creating new arcs with zero reduced cost, which can be used to push further flow from the origins to the destinations. If no such arcs can be created the problem is declared unfeasible. Much care is devoted to stop the max-flow computation as soon as a proof that the set of potentials is not optimal, in order to reach as soon as possible a dual optimal solution, and to re-use all available information to "warm start" the max-flow computation after a change in the potentials.

Warning

The original code has been written for integer data only. By properly setting the flow and cost tolerances [see `SetEps****()` in [MCFClass.h](#)] we have always been able to solve any MCF that we could throw at the solver, but in principle this kind of algorithm may fail to converge with nonintegral data, so consider yourselves warned.

A private type `SIndex` is defined which is intended to hold arc and node indices "with a sign", used to represent orientation. This has to be "in sync" with `Index`, in the sense that for every unsigned index value in `Index`, the two signed values should be feasible in `SIndex`. In other words, either `Index` is not using at least half of its feasible values, or `SIndex` has to be a "bigger" data type than `Index`. The default value for `SIndex` is `int`.

Version

1.80

Date

30 - 12 - 2009

Author**(original FORTRAN code)**

Dimitri P. Bertsekas

M.I.T.

(C++ porting and polishing)

Antonio Frangioni

Operations Research Group

Dipartimento di Informatica

Universita' di Pisa

(C++ porting and polishing)

Claudio Gentile

Istituto di Analisi di Sistemi e Informatica

Consiglio Nazionale delle Ricerche

Viale Manzoni, 30 Roma

Copyright &copy 1996 - 2009 by Antonio Frangioni